Assignment 2: Benchmarking

NSM vs. DSM and Isolation Levels

Special Topics in Computer Science: Big Data Systems CS-UH 3260 Spring 2023 MAX 10 points

<u>This assignment contains two parts.</u> The first focuses on analyzing the performance of two popular DBMSs with different storage models (DSM vs NSM). The second part involves simulating concurrency control issues.

PART I: In this assignment, you will evaluate two popular DBMSs featuring different storage models (DSM vs NSM), in addition to other optimization technologies. The performance evaluation will be based on real-world data obtained from the US Bureau of Transportation Statistics http://www.transtats.bts.gov/, which provides information on domestic flights within the United States.

1. System installation (not graded):

- Install MySQL, if not already installed, and create a database called "flight"
- Install MonetDB and follow instructions in lab3 monet.readme
 - note: call your database called "flight"

2. Database initialization and loading:

Next, you are going to load data from BTS collected for the years 2021 and 2022 (~7GB). The data has been preprocessed and is ready for you to create and populate a table called "ontime" with over 100 columns.

→ The data and the helper scripts are available on Drive.

A. Schema creation

- a. time mysql -uroot -p -D flight --local-infile=1 < air_ddl.sql
- b. time mclient -u monetdb -d flight < air_ddl.sql

B. Data Loading

- a. time mysql -uroot -p -D flight --local-infile=1 < load_mysql.sql
- b. time mclient -u monetdb -d flight < load_monetdb.sql (Note: put the full path to the data in the load sql script.)

@Deliverables:

- Report loading times. (0.5 point)
- Get the table status in both system: (0.5 point)
 - MySQL
 - SHOW TABLE STATUS;
 - MonetDB
 - select * from ontime limit 1;
 - select count(*) as c from ontime;

3. Benchmarking:

- A. Compare the execution time of MySQL and MonetDB to perform the SQL Queries¹ listed on the following page.
 - Advise: By default, the MySQL client provides the execution time of queries. However, in MonetDB, you need to activate the timer when starting the client using the following command:
 - o mclient -u monetdb -d flight --timer=performance

©Deliverables:

- Create histograms comparing the execution times of all five queries in MonetDB and MySQL, and analyze the performance of the systems based on their respective data storage models. (4 points)
- Propose ways to improve the performance of the slower system on one of the queries. (1point)
- B. <u>Automated Benchmarking (optional)</u>. As a stretch goal, create a script to automate the execution of all queries, the randomization of the variables, and the collection of the execution times.

©Deliverable:

Scripts in your favorite language (1 point <i>)

¹ The Queries are a subset of the Percona's Benchmark

Q1: Count records

SELECT count(*) form ontime;

Q2: Average monthly flights

SELECT avg(c1) FROM (
 SELECT YearD, MonthD, count(*) AS c1
 FROM ontime
 GROUP BY YearD, MonthD
) as tmp;

Q3: The number of flights per day from January to June in 2021

SELECT DayOfWeek, count(*) AS c FROM ontime WHERE YearD=2021 AND MonthD BETWEEN 1 AND 6 GROUP BY DayOfWeek ORDER BY c DESC;

Q4: The number of delays by carrier for 2022

SELECT Carrier, count(*)
FROM ontime
WHERE DepDelay>10 AND YearD=2022
GROUP BY Carrier
ORDER BY count(*) DESC;

Q5: Most popular destination by count of direct connected flights

SELECT

DestCityName AS destination_city, count(DISTINCT OriginCityName) AS num_origins_of_flights FROM ontime WHERE YearD BETWEEN 2021 AND 2021 GROUP BY DestCityName ORDER BY 2 DESC

LIMIT 10;

PART II: In this exercise, you will simulate issues related to concurrency control by modifying the default isolation level of MySQL

In your MySQL instance, create a new database and a simple relation. You may use the following code, or any other schema of your choice:

- CREATE DATABASE transaction_iso;
- USE transaction_iso;
- CREATE TABLE t (a INT NOT NULL, b INT, c INT);
- INSERT INTO t VALUES (1,2,3),(2, 4, 5),(3,12,10);

Then, using two mysql clients, simulate the following scenarios.

- Phantom read.
- Unrepeatable read.
- Dirty read.
- Lost update.

@Deliverable:

 Document the transactions that you performed to simulate each concurrency problem (1 point each (total 4pts))

Hints:

- Note your current isolation level using:
 - SHOW VARIABLES WHERE Variable_name LIKE "%_isolation";
- Make sure you are setting the sessions to the same isolation level. For example, to set the session into read uncommitted mode use:
 - SET SESSION TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
- Transactions are created, committed, and aborted as follows:
 - START TRANSACTION;
 - 0
 - COMMIT;
 - o ROLLBACK;
- These are the isolation levels in mysql:
 - READ UNCOMMITTED
 - READ COMMITTED
 - o REPEATABLE READ
 - o SERIALIZABLE