

Big Data Systems

Djellel Difallah

Spring 2023

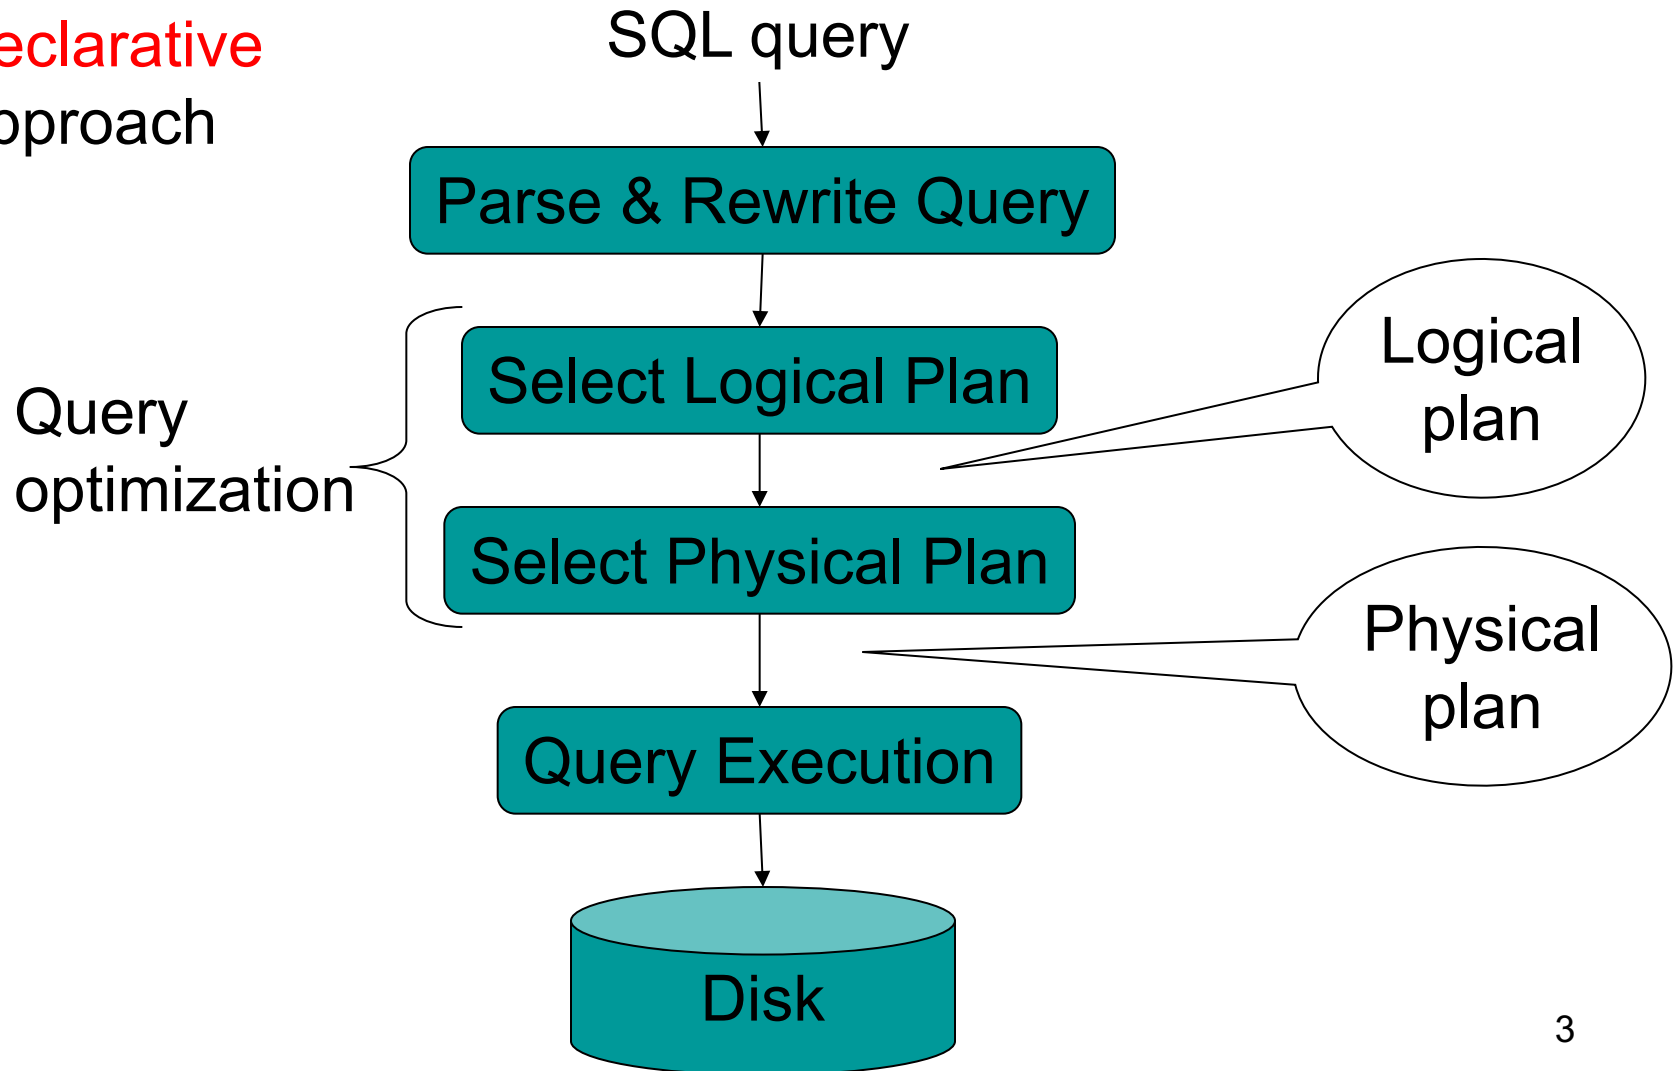
Lecture 3 – Query Execution

Outline

- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

Query Evaluation Steps

Declarative
Approach



Example Query

Supplier(sno, sname, scity, sstate)

Part(pno, pname, psize, pcolor)

Supply(sno, pno, price)

- Find the names of all suppliers in Seattle who supply part number 2

```
SELECT S.sname
```

```
FROM Supplier S, Supplies U
```

```
WHERE S.scity='Seattle' AND S.sstate='WA'
```

```
AND S.sno = U.sno
```

```
AND U.pno = 2;
```

Steps in Query Evaluation

- **Step 0: admission control**
 - User connects to the db with username, password
 - User sends query in text format
- **Step 1: Query parsing**
 - Parses query into an internal format
 - Performs various checks using catalog
 - Correctness, authorization, integrity constraints
- **Step 2: Query rewrite**
 - View rewriting, flattening, etc.

Continue with Query Evaluation


- **Step 3: Query optimization**
 - Find an efficient query plan for executing the query
- **A query plan is**
 - **Logical query plan**: an extended relational algebra tree
 - **Physical query plan**: with additional annotations at each node
 - Access method to use for each relation
 - Implementation to use for each relational operator

Extended Algebra Operators

- Union \cup , intersection \cap , difference $-$
- Selection σ
- Projection π
- Join \bowtie
- Duplicate elimination δ
- Grouping and aggregation γ
- Sorting τ
- Rename ρ

Logical Query Plan

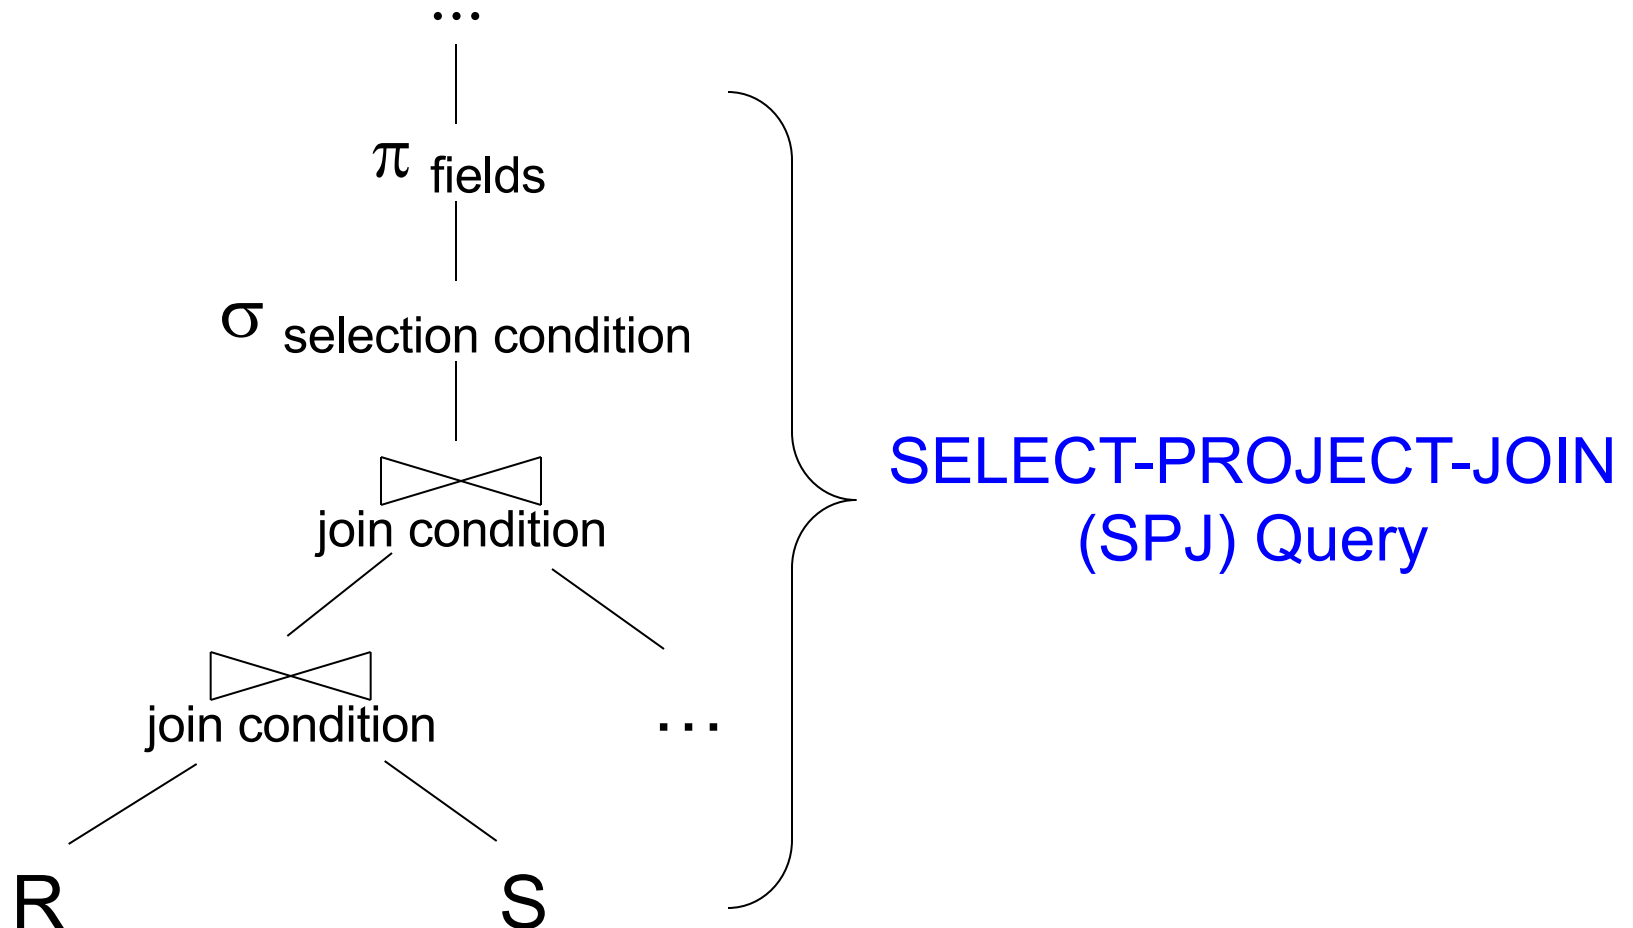
```
SELECT S.sname
FROM Supplier S, Supplies U
WHERE S.scity='Seattle'
AND S.sstate='WA'
AND S.sno = U.sno
AND U.pno = 2;
```



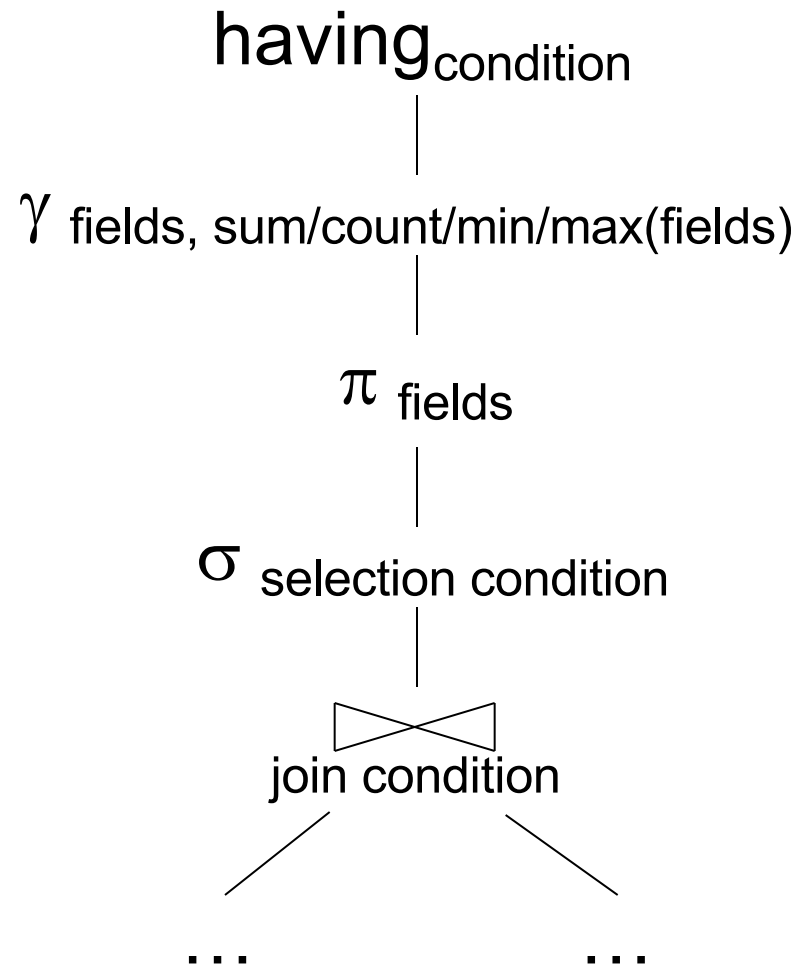
Query Block

- Most optimizers operate on individual query blocks
- A query block is a SQL query with **no nesting**
 - **Exactly one**
 - SELECT clause
 - FROM clause
 - **At most one**
 - WHERE clause
 - GROUP BY clause
 - HAVING clause

Typical Plan for Block (1/2)



Typical Plan For Block (2/2)



Physical Query Plan

- Logical query plan with extra annotations
- **Access path selection** for each relation
 - Use a file scan or use an index
- **Implementation choice** for each operator
- **Scheduling decisions** for operators

Physical Query Plan

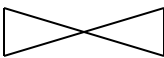
(On the fly)

π sname

(On the fly)

σ sscity='Seattle' \wedge sstate='WA' \wedge pno=2

(Nested loop)


sno = sno

Suppliers
(File scan)

Supplies
(File scan)

Final Step in Query Processing

- **Step 4: Query execution**
 - How to **synchronize operators**?
 - How to **pass data between operators**?
- Standard approach:
 - Record Iterator
 - Pipelined execution or Intermediate result materialization

Pipelined Execution

- Applies parent operator to tuples directly as they are produced by child operators
- Benefits
 - No operator synchronization issues
 - Saves cost of writing intermediate data to disk
 - Saves cost of reading intermediate data from disk
 - Good resource utilizations on single processor
- This approach is used whenever possible

Intermediate Tuple Materialization

- Writes the results of an operator to an intermediate table on **disk**
- Necessary for some operator implementations
 - E.g., When operator needs to examine the same tuples multiple times

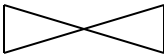
Pipelined Execution

(On the fly)

π_{sname}

(On the fly) $\sigma_{\text{sscity}='Seattle' \wedge \text{sstate}='WA' \wedge \text{pno}=2}$

(Nested loop)


 $\text{sno} = \text{sno}$

Suppliers
(File scan)


Supplies
(File scan)

Intermediate Tuple Materialization

(On the fly)

π_{sname}

(Sort-merge join)



(Scan: write to T1)

$\sigma_{\text{sscity}='Seattle' \wedge \text{sstate}='WA'}$

Suppliers
(File scan)

$\text{sno} = \text{sno}$

(Scan: write to T2)

$\sigma_{\text{pno}=2}$

Supplies
(File scan)

Outline

- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

Why Learn About Op Algos?

- Relevant for other Bid data systems
- Different systems implement different subsets of these algorithms
- Good algorithms can greatly improve performance

Cost Parameters

- In database systems the data is on disk
- **Cost = total number of I/Os**
 - More complex models possible (which ones?)
 - Different models for main-memory / distributed DBMSs (which ones?)
- Parameters:
 - $B(R)$ = # of blocks (i.e., pages) for relation R
 - $T(R)$ = # of tuples in relation R
 - $V(R, a)$ = # of distinct values of attribute a

Cost

- Cost of an operation = number of disk I/Os to
 - read the data
 - compute the result
- Cost of writing the result to disk is *not included*
 - Need to count it separately when applicable

Cost Parameters

- Clustered relation R:
 - $B(R) \approx T(R) / \text{blockSize}$
- Unclustered relation R:
 - Worst case: If the system uses shared blocks
 - $B(R) \approx T(R)$
- When **a** is a primary key, $V(R,a) = T(R)$
- When **a** is not a primary key, $V(R,a) = ?$

Cost of Scanning a Table

- Scanning a table:
 - Reading all its records.
- Clustered relation:
 - Result may be unsorted: $B(R)$
 - Result needs to be sorted: $3B(R)$ (see later)
- Unclustered relation
 - Unsorted: $T(R)$
 - Sorted: $T(R) + 2B(R)$

Outline

- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

One-pass Algorithms

Selection $\sigma(R)$, projection $\Pi(R)$

- Both are ***tuple-at-a-time*** operators
- Cost: $B(R)$, the cost of scanning a clustered relation



Join Algorithms

- Logical operator:
 - $\text{Product}(\text{pname}, \text{cname}) \bowtie \text{Company}(\text{cname}, \text{city})$
- Physical operators for the join:
 - **Hash join**
 - **Nested loop join**
 - **Sort-merge join**

Hash Join

Hash join: $R \bowtie S$

- Scan R, build buckets in main memory (M)
- Then scan S and join
- Cost: $B(R) + B(S)$
- **One pass algorithm** when one of the relations holds in memory!
 - That is: $B(R) \leq M$

Nested Loop Joins

- Tuple-based nested loop $R \bowtie S$
- R is the outer relation, S is the inner relation

```
for each tuple r in R do  
  for each tuple s in S do  
    if r and s join then output (r,s)
```

- Cost: $B(R) + T(R) B(S)$ when S is clustered
- Cost: $B(R) + T(R) T(S)$ when S is unclustered

Page-at-a-time Refinement

```
for each page of tuples r in R do  
  for each page of tuples s in S do  
    for all pairs of tuples  
      if r and s join then output (r,s)
```

- Cost: $B(R) + B(R)B(S)$ if S is clustered
- Cost: $B(R) + B(R)T(S)$ if S is unclustered

Nested Loop Joins

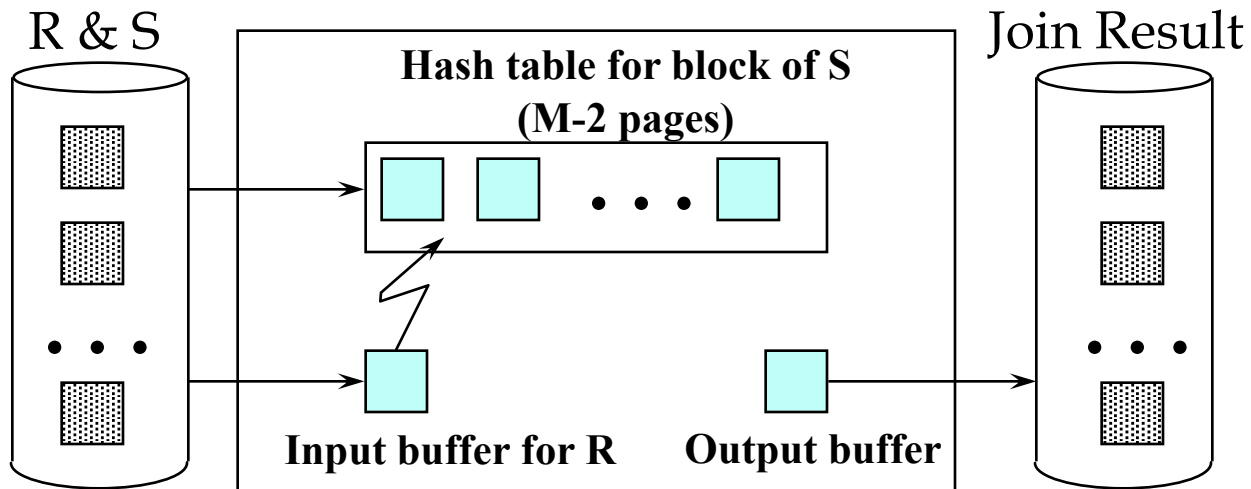
- We can be cleverer
- How would you compute the join in the following cases ?
What is the cost ?
 - $B(R) = 1000, B(S) = 2, M = 4$
 - $B(R) = 1000, B(S) = 3, M = 4$
 - $B(R) = 1000, B(S) = 6, M = 4$

Block Nested Loop Joins

- Block Nested Loop Join
- Group of $(M-2)$ pages of S is called a “block”

```
for each  $(M-2)$  pages  $ps$  of  $S$  do  
  for each page  $pr$  of  $R$  do  
    for each tuple  $s$  in  $ps$   
      for each tuple  $r$  in  $pr$  do  
        if “ $r$  and  $s$  join” then output( $r,s$ )
```


Block Nested Loop Joins



Block Nested Loop Joins

- Cost of block-based nested loop join
 - Read S once: cost $B(S)$
 - Outer loop runs $B(S)/(M-2)$ times, and each time need to read R: costs $B(S)B(R)/(M-2)$
 - Total cost: $B(S) + B(S)B(R)/(M-2)$
- Notice: it is better to iterate over the smaller relation first

Sort-Merge Join

Sort-merge join: $R \bowtie S$

- Scan R and sort in main memory
- Scan S and sort in main memory
- Merge R and S
- Cost: $B(R) + B(S)$
- One pass algorithm when $B(S) + B(R) \leq M$
- Typically, this is NOT a one pass algorithm

More One-pass Algorithms

Duplicate elimination $\delta(R)$

- Need to keep tuples in memory
- When new tuple arrives, need to compare it with previously seen tuples
- Balanced search tree or hash table
- Cost: $B(R)$
- Assumption: $B(\delta(R)) \leq M$

Even More One-pass Algorithms

Grouping:

Product(name, department, quantity)

$\gamma_{\text{department, sum(quantity)}}(\text{Product})$

How can we compute this in main memory ?

Even More One-pass Algorithms

- Grouping: $\gamma_{\text{department, sum(quantity)}}(R)$
- Need to store all departments in memory
- Also store the sum(quantity) for each department
- Balanced search tree or hash table
- Cost: $B(R)$
- Assumption: number of depts fits in memory

Outline


- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

Two-Pass Algorithms

- What if data does not fit in memory?
- Need to process it in multiple passes
- Two key techniques
 - Hashing
 - Sorting
 - Write back to disk!

Two Pass Algorithms Based on Hashing

- Idea: partition a relation R into buckets, on disk
- Each bucket has size approx. $B(R)/M$



- Does each bucket fit in main memory ?
 - Yes if $B(R)/M \leq M$, i.e. $B(R) \leq M^2$

Hash Based Algorithms for δ

- Recall: $\delta(R)$ = duplicate elimination
- Step 1. Partition R into buckets
- Step 2. Apply δ to each bucket
- Cost: $3B(R)$
- Assumption: $B(R) \leq M^2$


Partitioned (Grace) Hash Join

$R \bowtie S$

- Step 1:
 - Hash S into M-1 buckets
 - Send all buckets to disk
- Step 2
 - Hash R into M-1 buckets
 - Send all buckets to disk
- Step 3
 - Join every pair of buckets

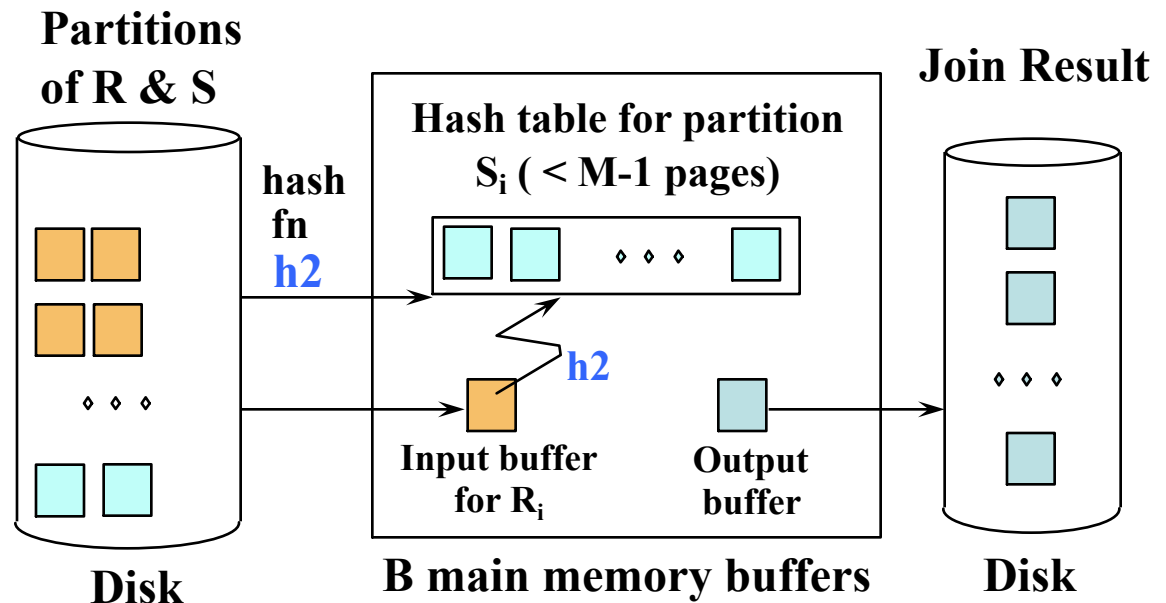
Partitioned Hash Join

- Partition both relations using hash function **h**
- R tuples in partition i will only match S tuples in partition i .



Partitioned Hash Join

- Read in partition of R, hash it using $h_2 (\neq h)$
 - **Build phase**
- Scan matching partition of S, search for matches
 - **Probe phase**



Partitioned Hash Join


- Cost: $3B(R) + 3B(S)$
- Assumption: $B(R)$ and $B(S) \leq M^2$

External Sorting

- Problem: Sort a file of size B with memory M
- E.g., ORDER BY in SQL queries
- Will discuss only 2-pass sorting, for when $B < M^2$

External Merge-Sort: Step 1


- Phase one: load M pages in memory, sort



Runs of length M pages

External Merge-Sort: Step 2

- Merge $M - 1$ runs into a new run
- Output: run of length $M (M - 1) \approx M^2$



If $B(R) \leq M^2$ then we are done
(slightly more complex otherwise)

External Merge-Sort

- Cost:
 - Read+write+read = $3B(R)$
 - Assumption: $B(R) \leq M^2$
- Other considerations
 - In general, a lot of optimizations are possible

Two-Pass Algorithms Based on Sorting

Duplicate elimination $\delta(R)$

- Trivial idea: sort first, then eliminate duplicates
- Step 1: sort chunks of size M , write
 - cost $2B(R)$
- Step 2: merge $M-1$ runs, but include each tuple only once
 - cost $B(R)$
- Total cost: $3B(R)$, Assumption: $B(R) \leq M^2$

Two-Pass Algorithms Based on Sorting

Grouping: $\gamma_{a, \text{sum}(b)}(R)$

- Same as before: sort, then compute the $\text{sum}(b)$ for each group of a 's
- Total cost: $3B(R)$
- Assumption: $B(R) \leq M^2$

Two-Pass Algorithms Based on Sorting

Join $R \bowtie S$

- Start by sorting both R and S on the join attribute:
 - Cost: $4B(R)+4B(S)$ (because need to write to disk)
- Read both relations in sorted order, match tuples
 - Cost: $B(R)+B(S)$
- Total cost: $5B(R)+5B(S)$
- Assumption: $B(R) \leq M^2$, $B(S) \leq M^2$

Two-Pass Algorithms Based on Sorting

Join $R \bowtie S$

- Also, if $B(R) + B(S) \leq M^2$ we can compute the join during the merge phase
- Total cost: $3B(R) + 3B(S)$

Outline

- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

Review: Access Methods

- **Heap file**
 - Scan tuples one at the time
- **Hash-based index**
 - Efficient selection on equality predicates
- **Tree-based index**
 - Efficient selection on equality or range predicates

Index Based Selection

- Selection on equality: $\sigma_{a=v}(R)$
- $V(R, a) = \text{\textit{\# of distinct values of attribute a}}$
- Clustered index on a : $\text{cost} \sim B(R)/V(R,a)$
- Unclustered index on a : $\text{cost} \sim T(R)/V(R,a)$
- Note: we ignored the I/O cost for the index pages

Index Based Selection

- Example:
 - $B(R) = 2000$
 - $T(R) = 100,000$
 - $V(R, a) = 20$
 - Compute the cost of $\sigma_{a=v}(R)$
- Index based selection
 - If index is clustered: $B(R)/V(R,a) = 100$
 - If index is unclustered: $T(R)/V(R,a) = 5,000$
- Note
 - Don't build unclustered indexes when $V(R,a)$ is small.


Index Nested Loop Join

$R \bowtie S$ (Natural Join)

- Assume S has an index on the join attribute
- Iterate over R, for each tuple fetch corresponding tuple(s) from S
- Cost:
 - Assuming R is clustered
 - If index on S is clustered: $B(R) + T(R)B(S)/V(S,a)$
 - If index on S is unclustered: $B(R) + T(R)T(S)/V(S,a)$

Index Based Join

- Assume both R and S have a sorted index (B+tree) on the join attribute
- Perform a merge join (Zig-Zag join)
- Cost: $B(R) + B(S)$



Outline

- **Steps involved in processing a query**
 - Logical query plan
 - Physical query plan
 - Query execution overview
- **Operator implementations**
 - One pass algorithms
 - Two-pass algorithms
 - Index-based algorithms

Summary of Query Execution

- For each logical query plan
 - There exist many physical query plans
 - Each plan has a different cost
 - Cost depends on the data
- Additionally, for each query
 - There exist several logical plans
- Query optimization factors in the cost of several logical plans by using the operators and data estimates (very fast)