

Big Data Systems

Djellel Difallah

Spring 2023

Lecture 6 – Document Stores

MongoDB

Different Types of NoSQL

Taxonomy based on data models [Rick Cattell, SIGMOD Record 2011]

- Key-value stores
 - Value is arbitrary (think of it as a scalable hashmap)
 - e.g., Project Voldemort, Memcached, **Dynamo**
- Column Family (a.k.a Extensible Record Stores)
 - Some schema structure
 - e.g., HBase, **Cassandra**, PNUTS
- **Document stores**
 - Complex data structure (JSON seems to be the standard but can be XML)
 - e.g., SimpleDB, CouchDB, **MongoDB**

Agenda

- Document Stores Overview
- MongoDB Introduction
- Data Model
- Query Operations
- Consistency and Replication
- Storage Engine
- MongoDB vs Cassandra

Why Document Store?

- There is a big gap between the application and the data representation.
- Central concept: the Document
 - Similar, in some ways, to records or rows in relational databases, but are less rigid
 - They do not required to adhere to a strict schema, nor will they have all the same sections, slots, parts, or keys.
- Designed for storing, retrieving, and managing document-oriented information such as semi-structured data, e.g.,
 - JSON, XML, YAML, CSV, etc.

Why Document Store? (cont.)

- JSON is popular for data exchange (XML less nowadays)
- Data stored in document DB can be used directly
- Databases often store objects from memory
 - Using RDBMS, we must do Object Relational Mapping (ORM)
 - ORM is relatively demanding
 - JSON is much closer to structure of memory objects
 - It was originally for JavaScript objects
 - Object Document Mapping (ODM) is faster

Document Store Features

- Semi-structured data => optional schema
- Ad hoc queries (search by field, range, etc.)
- Often provide:
 - RESTful API
 - Data aggregation API (e.g., through Map/Reduce)
 - Full text search
 - Secondary indexes
 - Automatic sharding (scale writes)
- ... but are typically missing:
 - Explicit locks
 - Strong consistency guarantees
 - Transactions
 - Joins
 - *Literally* No SQL

Flexible Schemas

Country	France
Region	Île-de-France
Department	Paris
Intercommunality	Métropole du Grand Paris
Subdivisions	20 arrondissements
Government	
• Mayor (2020–2026)	Anne Hidalgo ^[1] (PS)
Area ¹	105.4 km ² (40.7 sq mi)
• Urban (2020)	2,853.5 km ² (1,101.7 sq mi)
• Metro (2020)	18,940.7 km ² (7,313.0 sq mi)
Population (Jan. 2019) ^[2]	2,165,423
• Density	21,000/km ² (53,000/sq mi)
• Urban (2019) ^[3]	10,858,852
• Urban density	3,800/km ² (9,900/sq mi)
• Metro (Jan. 2017) ^[4]	13,024,518
• Metro density	690/km ² (1,800/sq mi)
Demonym(s)	Parisian(s) ^(en) <i>Parisien(s)</i> (masc.), <i>Parisienne(s)</i> (fem.) ^(fr) <i>Parigot(s)</i> (masc.), "Parigote(s)" (fem.) (fr, colloquial)
Time zone	UTC+01:00 (CET)
• Summer (DST)	UTC+02:00 (CEST)
INSEE/Postal code	75056 ^(fr) /75001-75020, 75116
Elevation	28–131 m (92–430 ft) (avg. 78 m or 256 ft)
Website	www.paris.fr ^(fr)

Abu Dhabi

Country	United Arab Emirates
Emirate	Abu Dhabi
Municipal region	Central Capital District ^[1]
Government	
• Type	Municipality
• Body	Abu Dhabi City Municipality
• Director-General of City Municipality	Saif Badr al-Qubaisi
Area	
• Total	972 km ² (375 sq mi)
Elevation	27 m (89 ft)
Population (2021) ^{[2] [3]}	
• Total	1,512,000
• Density	1,600/km ² (4,000/sq mi)
Demonyms	Abu Dhabiyan, Dhabyani
Time zone	UTC+4 (UAE Standard Time)
GDP PPP	2014 estimate
Total	US\$ 178 billion ^[4]
Per capita	US\$ 61,000
Website	tamm.abudhabi ^(en)

Abu Dhabi

Country	 United States
Residence Act	1790
Organized	1801
Consolidated	1871
Home Rule Act	1973
Named for	George Washington, Christopher Columbus
Government	
• Mayor	Muriel Bowser (D)
• D.C. Council	List ^(en) ^(show)
• U.S. House	Eleanor Holmes Norton (D), Delegate (At-large)
Area	
• Federal capital city and federal district	68.35 sq mi (177.0 km ²)
• Land	61.126 sq mi (158.32 km ²)
• Water	7.224 sq mi (18.71 km ²)
Highest elevation	409 ft (125 m)
Lowest elevation	0 ft (0 m)
Population (2020) ^[2]	
• Federal capital city and federal district	689,545
• Estimate (2021) ^[2]	670,050
• Rank	23rd in the United States
• Density	11,280.71/sq mi (4,355.39/km ²)
• Urban ^[3]	5,174,759 (US: 8th)
• Urban density	3,997.5/sq mi (1,543.4/km ²)
• Metro ^[4]	6,385,162 (US: 6th)
Demonym	Washingtonian ^{[5][6]}
Time zone	UTC−5 (EST)
• Summer (DST)	UTC−4 (EDT)
ZIP Codes	20001–20098, 20201–20599, 56901–56999
Area code(s)	202, 771 (overlay) ^{[7][8]}
International airports	Dulles International Reagan National Baltimore/Washington
Commuter rail	MARC Train Virginia Railway Express
Rapid transit	Washington Metro
Website	dc.gov ^(en)

Washington D.C

Document Example

```
{
  FirstName:"Camilla",
  Address:"Saadiyat",
  Children:[
    { Name:"Mike",   Age:10},
    { Name:"Jennie", Age:8},
    { Name:"Samantha", Age:5},
    { Name:"Ines",   Age:2}
  ]
},
{
  FirstName:"Omar",
  Address:"Yas Island",
  Hobby:"sailing"
}
```

Can be serialized in CSV, JSON, XML, etc.

JSON

- JSON: Text-based open standard for data interchange
Serializing and transmitting structured data
 - JSON = JavaScript Object Notation Derived from JavaScript scripting language Uses conventions of the C-family of languages
 - Language independent (see www.json.org)
- BSON (Binary JSON) is a binary-encoded serialization format used to store and transfer data in a more compact and efficient manner compared to JSON.

NOSQL Document Stores

MONGODB



** Illustrative figures from mongodb.com*

Data Model

- Collections of documents (analogue of a table)
- BSON documents (attribute-value pairs with nesting and arrays)
- Documents can reference each other
 - Apps must issue follow-up queries to resolve the references
- Documents can be embedded in each other (i.e., nested)
 - But then we have to worry about documents getting very large

Data Model (cont.)

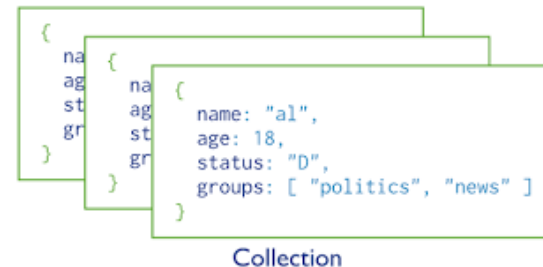
- Documents have flexible schema
 - Collections do not enforce specific data structure
- Related data in a single document structure
 - Documents can have subdocuments (in a field or array)
- Key decision of data modeling:
 - References vs. embedded documents
 - In other words: another trade-off
 - Locality of data
 - Relationships between data

Relational vs. MongoDB

Relational	MongoDB Model
Database	Database
Table	Collection
Tuple	Document (BSON)
Row_id	_id
Column	Field in the document

Each JSON document:

- belongs to a collection
- has a field `_id`
 - Unique within the collection
 - Primary Key
 - Immutable
 - Can be generated automatically



Operations

- Javascript API and Javascript shell
- Querying:Selection
 - A query targets a collection of documents
 - Selection queries on attribute values (including arrays)
`db.inventory.find({ type: "snacks" })`
 - Can also have conditions on embedded documents
 - Can also do projections, sort, limit and skip
- Querying:Aggregation
 - Aggregation pipelines `db.orders.distinct("cust_id")`
 - MapReduce API (JavaScript map/reduce functions)

Operation Examples

```
db.users.find(
  { age: { $gt: 18 } },
  { name: 1, address: 1 }
).limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

```
db.users.insertOne(
  {
    name: "sue",
    age: 26,
    status: "pending"
  }
)
```

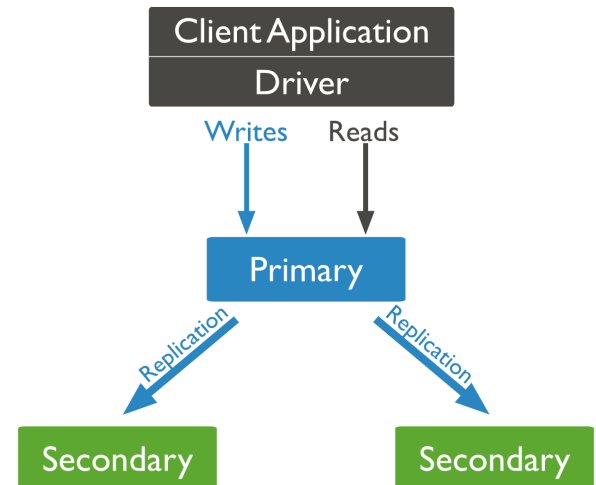
← collection
← field: value
← field: value
← field: value } document

```
db.users.updateMany(
  { age: { $lt: 18 } },
  { $set: { status: "reject" } }
)
```

← collection
← update filter
← update action

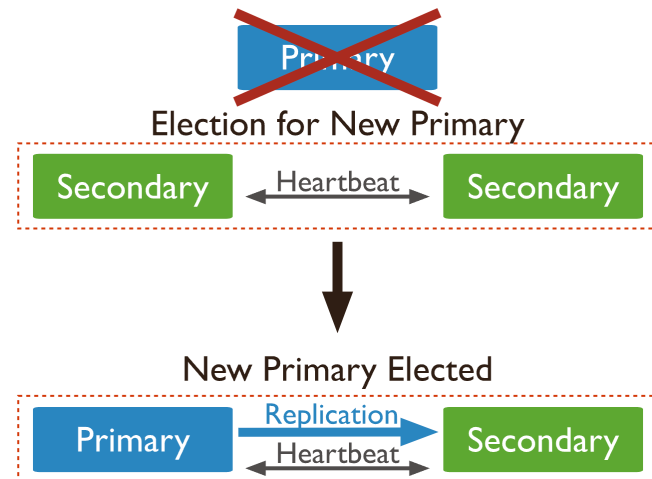
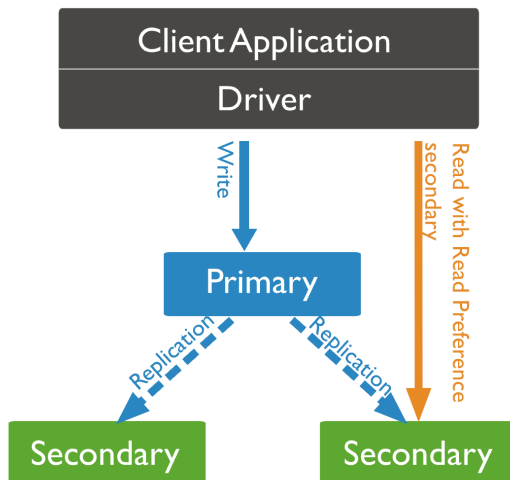
Replication in MongoDB

- MongoDB uses asynchronous replication for high availability.
- A **Replica Set** is a group of nodes that maintain the same data (max 50)
- **The primary node** is the master node that receives all write operations.
 - The primary records all changes to its data sets in its operation log or oplog (similar to a commit log)
- **Secondary nodes** replicate the primary's oplog and apply the operations to their data sets.
- **Strongly consistent system.** Once a write completes, any subsequent read will return the most recent value.



Replication in MongoDB

- **Read Preference:** MongoDB client can route some or all reads to the secondary members optionally, but writes must be sent to the primary.
- **Failover:** If the primary is unavailable, an eligible secondary will hold an **election** to elect itself the new primary.



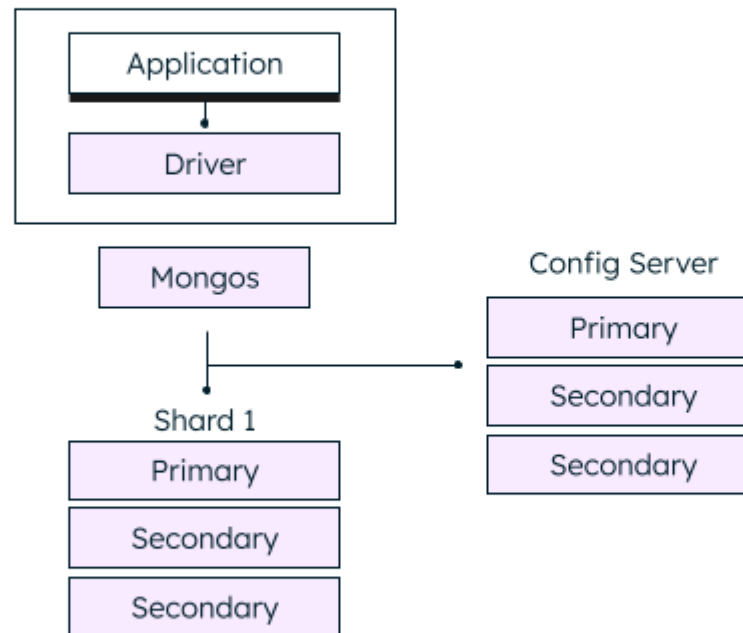
Replication and Consistency

- Consistency
 - Write operations are atomic at the document level
 - Operations that modify more than a single document in a collection still operate on one document at a time
- Replication
 - Master scheme with master performing all reads & writes by default
 - Achieves strong consistency by always going through the master
 - Eventual consistency by default when reading from replicas
 - Updates propagate to replicas asynchronously
 - But can be configured to use synchronous replication

Sharding in MongoDB

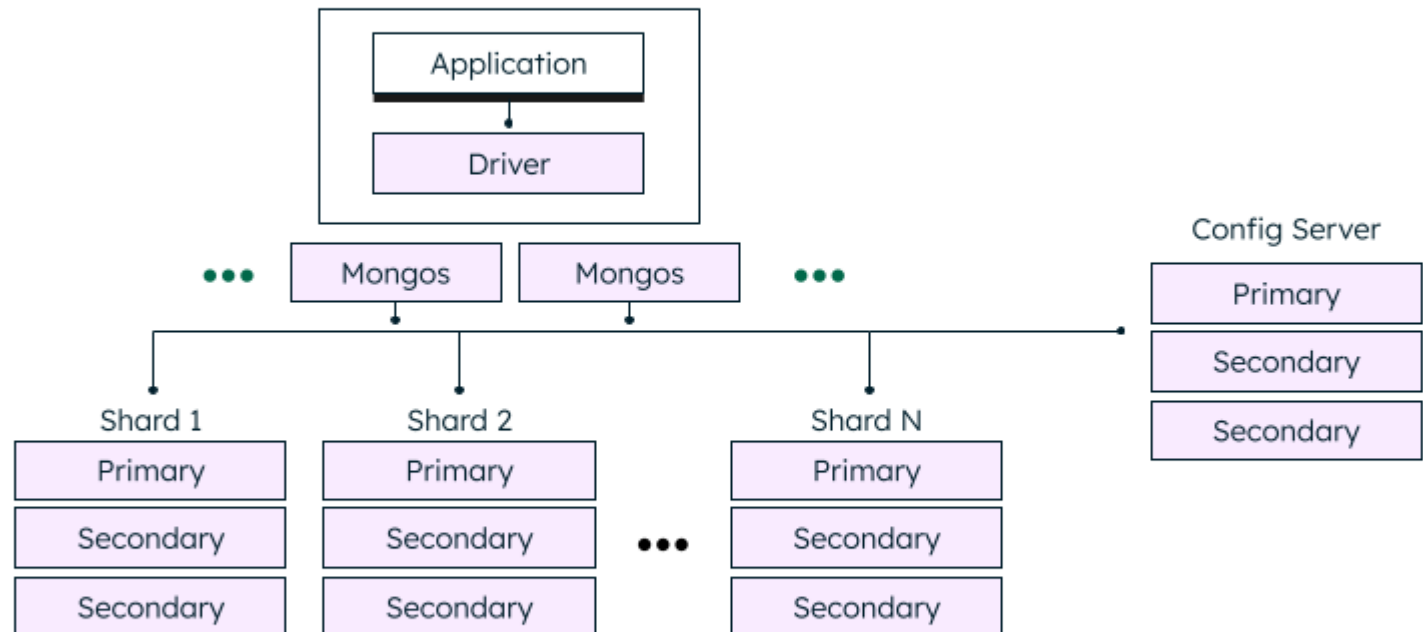
- Sharding? Partitioning? 🤔
 - Sharding is the concept of splitting a database's table (in mongo's case a collection) across multiple machines.
- MongoDB implements **sharding** using the following components:
 - **shard**: Each shard contains a subset of the sharded data. Each shard can be deployed as a replica set.
 - **mongos**: The mongos acts as a query router, providing an interface between client applications and the sharded cluster.
 - **config servers**: Config servers store metadata and configuration settings for the cluster. Deployed as replica set.
- User selects the “shard key”

MongoDB Architecture



Local Deployment

MongoDB Architecture



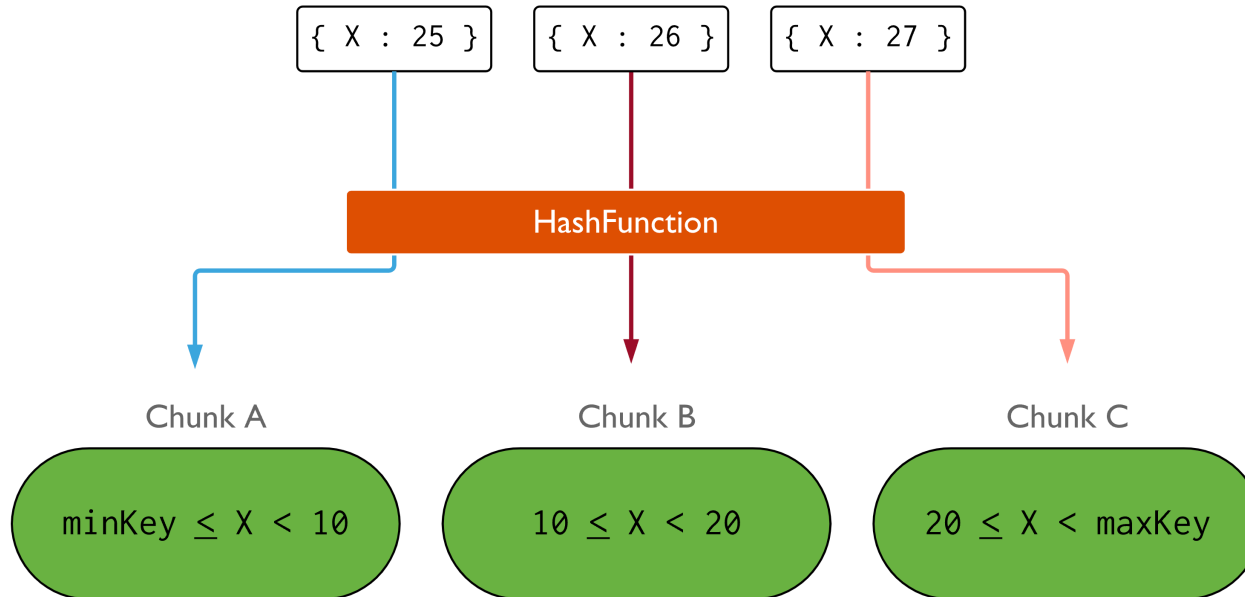
Cluster Deployment

Tuning Consistency

- Set read preference to secondary nodes.
 - Primary
 - Secondary
 - Nearest
 - Hedged reads
- Write Concern
 - { w: <value>, j: <written to journal>, wtimeout: <timeout> }
- Read Concern
 - “Local” (most recent data)
 - “Majority” (data has been replicated to a majority of secondaries)

Sharding Strategies in MongoDB

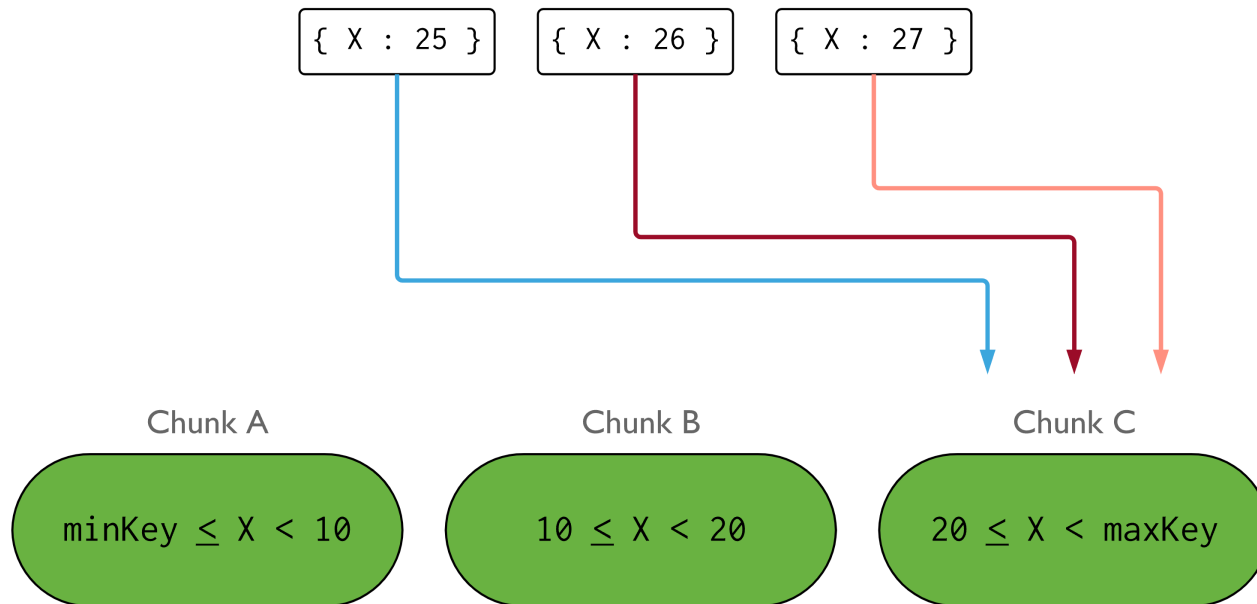
- Hash Sharding



```
sh.shardCollection("<database>.<collection>", { <shard key field> : "hashed" } )
```

Sharding Strategies in MongoDB

- Range Sharding



```
sh.shardCollection("<database>.<collection>", { <shard key field> : 1 } )
```


Storage Engines

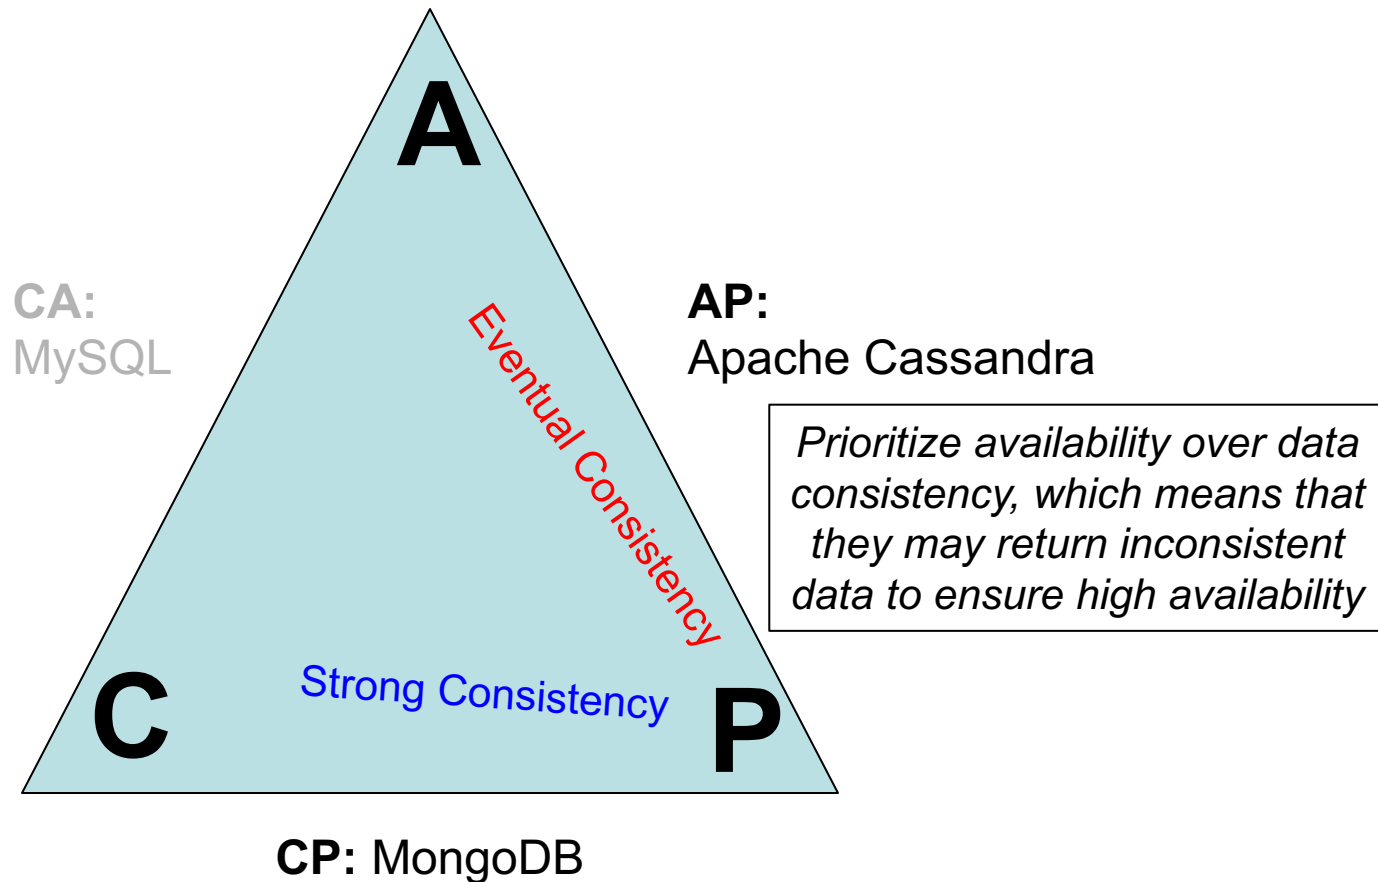
- MMAPv1 storage engine (historical)
 - Filesystem mmap: volume inserts, reads, and in-place updates.
- WiredTiger (starting version 3.2) which provides:
 - Document-level concurrency control for write operations
 - Optimistic concurrency control
 - Durability:
 - Multi-version concurrency control (MVCC)
 - Write ahead log (WAL), or Journal

Additional Concepts

- Indexing
 - MongoDB automatically indexes the `_id` field
 - Users can add indexes on other attributes (secondary index)
 - Geospatial Indexing
- Data modeling techniques: Document nesting vs. Referencing
- More features
 - Can insert document with a “Time To Live” to expire date
 - FIFO management of inserted docs and efficiently support operations that insert/read docs based on insertion order (Capped collections)

```
db.createCollection("recent", { capped : true, max : 10000 } )
```

NOSQL



To maintain data consistency, the system sacrifices availability by rejecting the request or cancelling it.

NoSQL - Takeaways

- Understand the rough limits of systems based on their architecture, data models, tradeoff
- MongoDB vs Cassandra
 - Data Model: MongoDB stores data in JSON-like format, while Cassandra uses a column-family (rows and columns)
 - Query: MongoDB has a rich query language and supports ad-hoc queries, while Cassandra has a limited query language and **requires data to be modeled based on the queries to be performed.**
 - Scalability: Both databases are horizontally scalable, but Cassandra is designed to handle massive amounts of data across multiple nodes more efficiently (decentralized).
 - Consistency: MongoDB provides strong consistency by default, while Cassandra provides eventual consistency by default. Both allow for **tunable consistency levels.**

Installation

- Install MongoDB
 - On MacOS: `brew install mongodb-community@6.0`
 - Start the client: `./mongosh`
- More info:
 - <http://cassandra.apache.org>