

# Big Data Systems

Djellel Difallah

Spring 2023

Lecture 7 – More NOSQL:  
Graph DBMS, Array DBMS

# Agenda

---

- Graph DBMS
  - Introduction
  - Storage and Indexing
  - Partitioning
- Array DBMS
  - Introduction
  - Physical Array Storage
  - Space filling curves

# Small vs Large Graphs

---

- Small graphs
  - Manage a collection of small graphs
  - Bioinformatics and Cheminformatics
  - Usually, data fit in memory (well studied)
- Large graphs
  - One large graph, aka “network”
  - Social network
  - Knowledge graphs
  - Active area of research


# Random vs Scale-free Graphs

---


- Random graphs
  - Node degree is constrained
- Scale-free graphs
  - Distribution of node degree follows a power law distribution
  - Most large graphs are scale-free
  - Small world phenomena & hubs
  - hard to partition

# Random vs Scale-free Graphs


**a.** POISSON



**b.**



**c.** POWER LAW





**d.**



# Graphs: a unique Data Model

- While the relational model is flexible to represent a large , specialized models tends to gives huge execution benefits



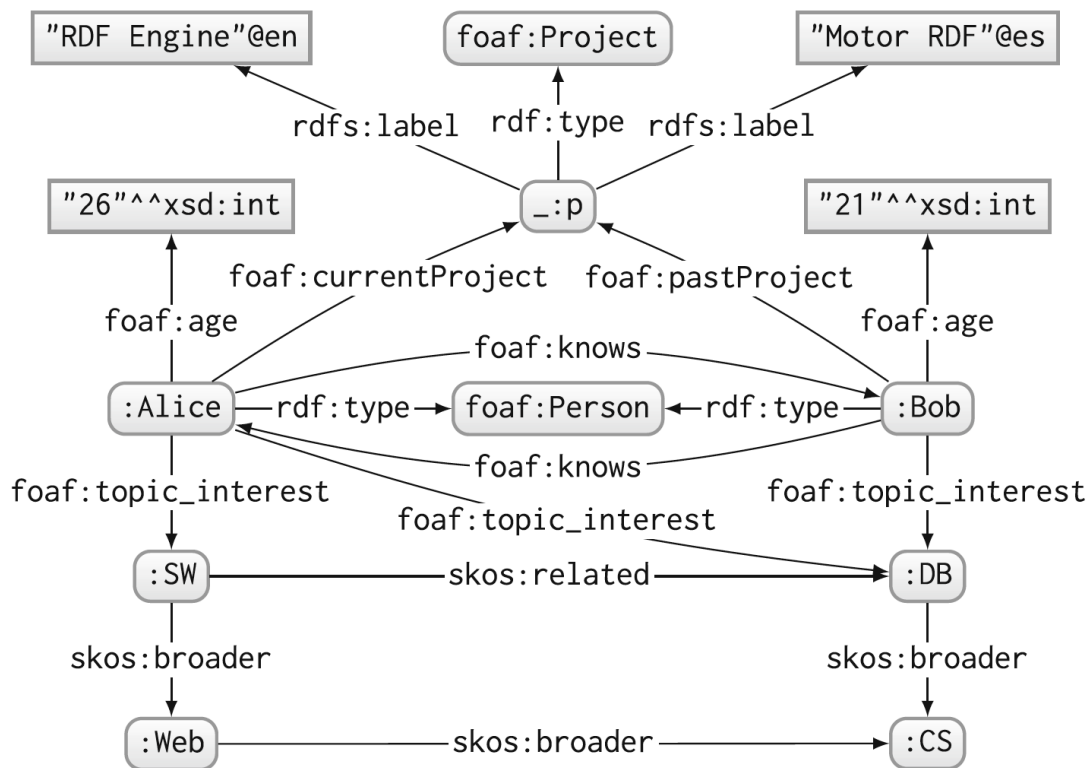
# Graph Database

---

- Optimized for the connections between records
  - Fast querying across records
- Transactional with CRUD operations (create/read/update/delete)
- Difference in the use case:
  - A relational database may tell you how many books your read last year
  - A graph database will tell you *which books you and the friends of your friends have read in common*
- Query Languages:
  - Many specialized and/or proprietary languages
  - SPARQL standard for RDF data

# Triple Pattern and Basic Graph Pattern

- Recall triples: s (subject), p (predicate), o (object).
  - The unit information in a property graph




Subject	Predicate	Object
:Alice	rdf:type	foaf:Person
:Alice	foaf:age	"26"^^xsd:int
:Alice	foaf:topic_interest	:DB
:Alice	foaf:topic_interest	:SW
:Alice	foaf:knows	:Bob
:Alice	foaf:currentProject	_:p
:Bob	rdf:type	foaf:Person
:Bob	foaf:age	"21"^^xsd:int
:Bob	foaf:topic_interest	:DB
:Bob	foaf:knows	:Alice
:Bob	foaf:pastProject	_:p
_:p	rdf:type	foaf:Project
_:p	rdfs:label	"RDF Engine"@en
_:p	rdfs:label	"Motor RDF"@es
:SW	skos:broader	:Web
:SW	skos:related	:DB
:Web	skos:broader	:CS
:DB	skos:broader	:CS



# Triple Pattern and Basic Graph Pattern

- **Triple patterns** are used to select sets of triples:
  - Example: `?y :rdf-type foaf:Person`
  - This is an (\*, p, o) pattern, ie., matches all subjects with the provided predicate and object
  - Binds all subjects to the variable `?y`
- A basic pattern is a set of triple patterns.
  - Returns triples that match all patterns (intersection)

```
SELECT * WHERE {  
  ?a a foaf:Person ; foaf:knows ?b ; foaf:topic_interest ?ia .  
  ?b a foaf:Person ; foaf:knows ?a ; foaf:topic_interest ?ib .  
}
```



?a	?b	?ia	?ib
:Alice	:Bob	:DB	:DB
:Alice	:Bob	:SW	:DB
:Bob	:Alice	:DB	:DB
:Bob	:Alice	:DB	:SW

---

# **GRAPH DATABASES**

## **STORAGE**

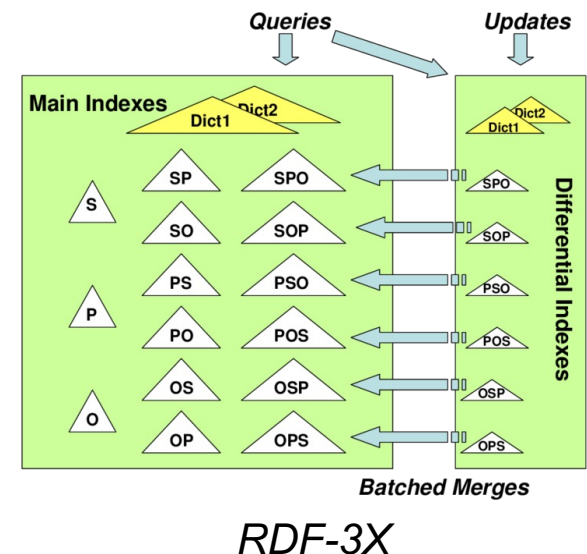
# Graph Data Structures

---

- Adjacency Matrix
  - Can use matrix operations / Linear algebra
  - Array stores
  - Inneficient storage (high sparsity)
- Adjacency lists
  - Very good for graph search
  - Document stores
  - Redundancy
- List of triples
  - Can use relational models (Triple Table, Property tables, Vertical partitioning)
  - Inner joins!

# Index-based Triple Stores

- SIX indexes (B+ trees):
  - All permutations of S, P, O
  - Single index scan
- Pros
  - Highly compressed
- Cons
  - Complex query execution engines
- Systems:
  - **RDF-3X**: “The RDF-3X engine for scalable management of RDF data”
  - **Hexastore**: “Sextuple Indexing for Semantic Web Data Management”

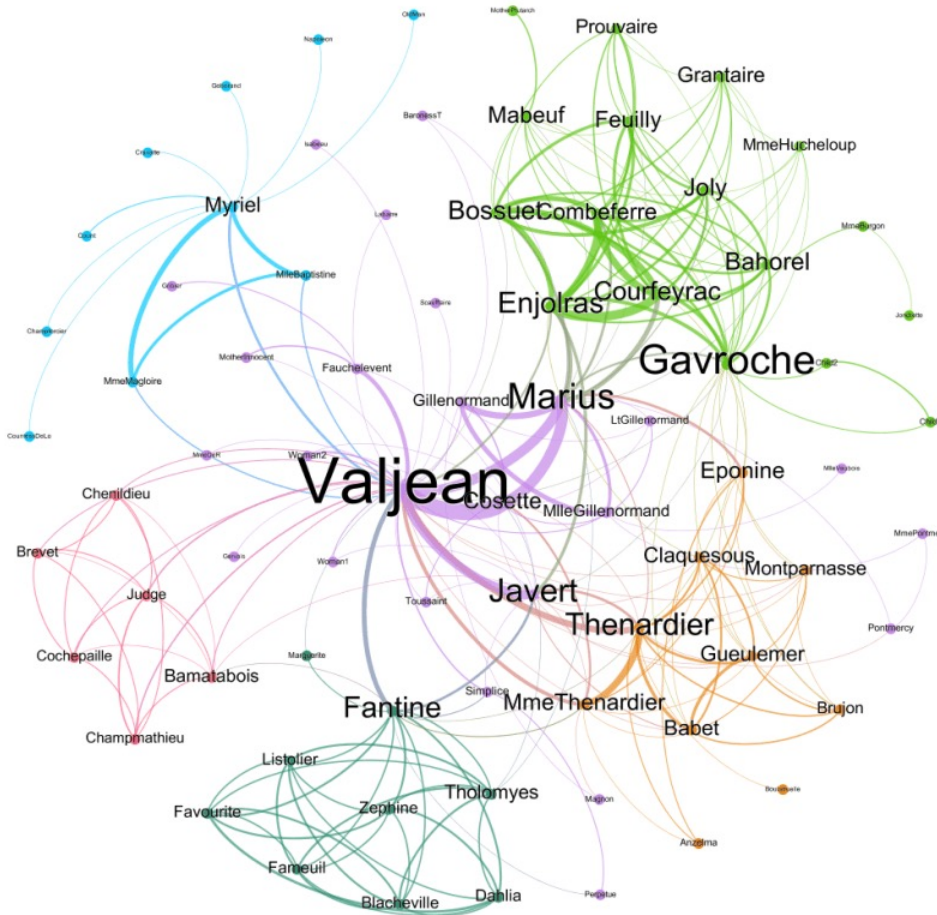


---

# **GRAPH DATABASES**


## **PARTITIONING**

# Partitioning




*Les Misérables* Character Co-occurrence Graph

# Adjacency Matrix



[Visualization Link](#)

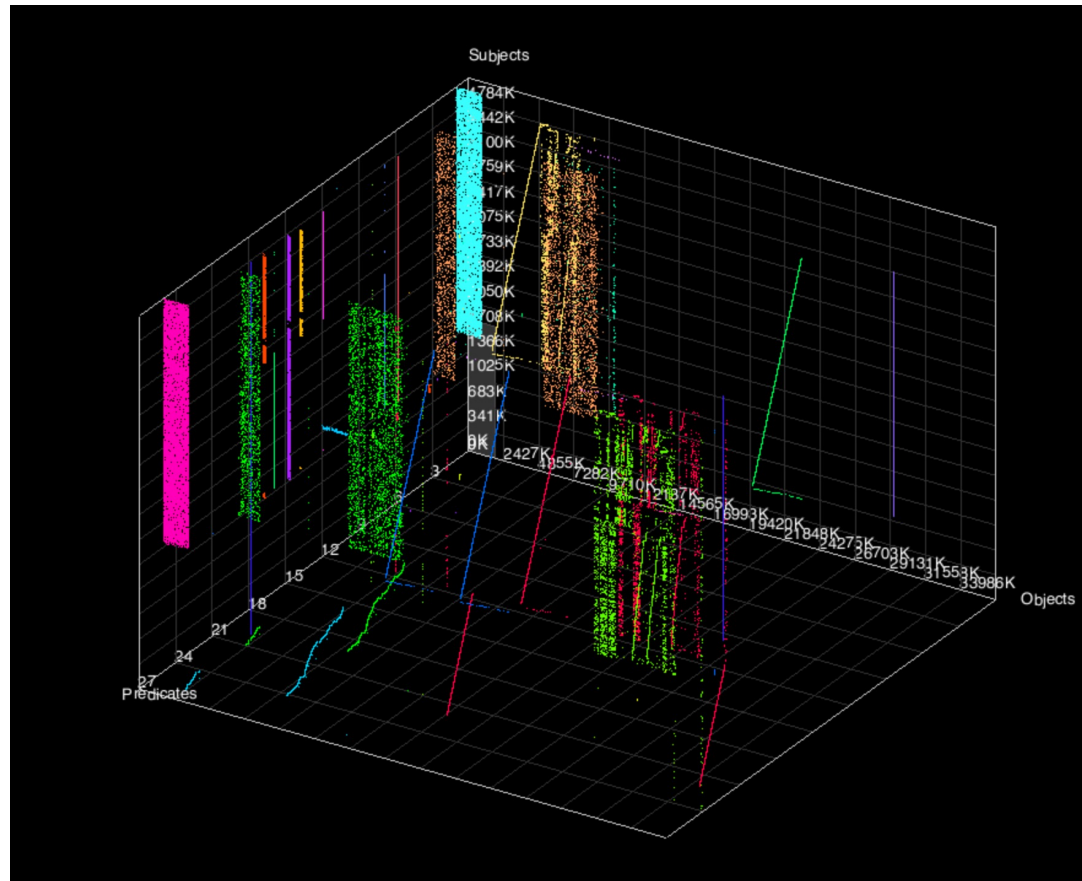
# Adjacency Matrix



[Visualization Link](#)



# 3D Matrix



RDF HDT

# Agenda

---

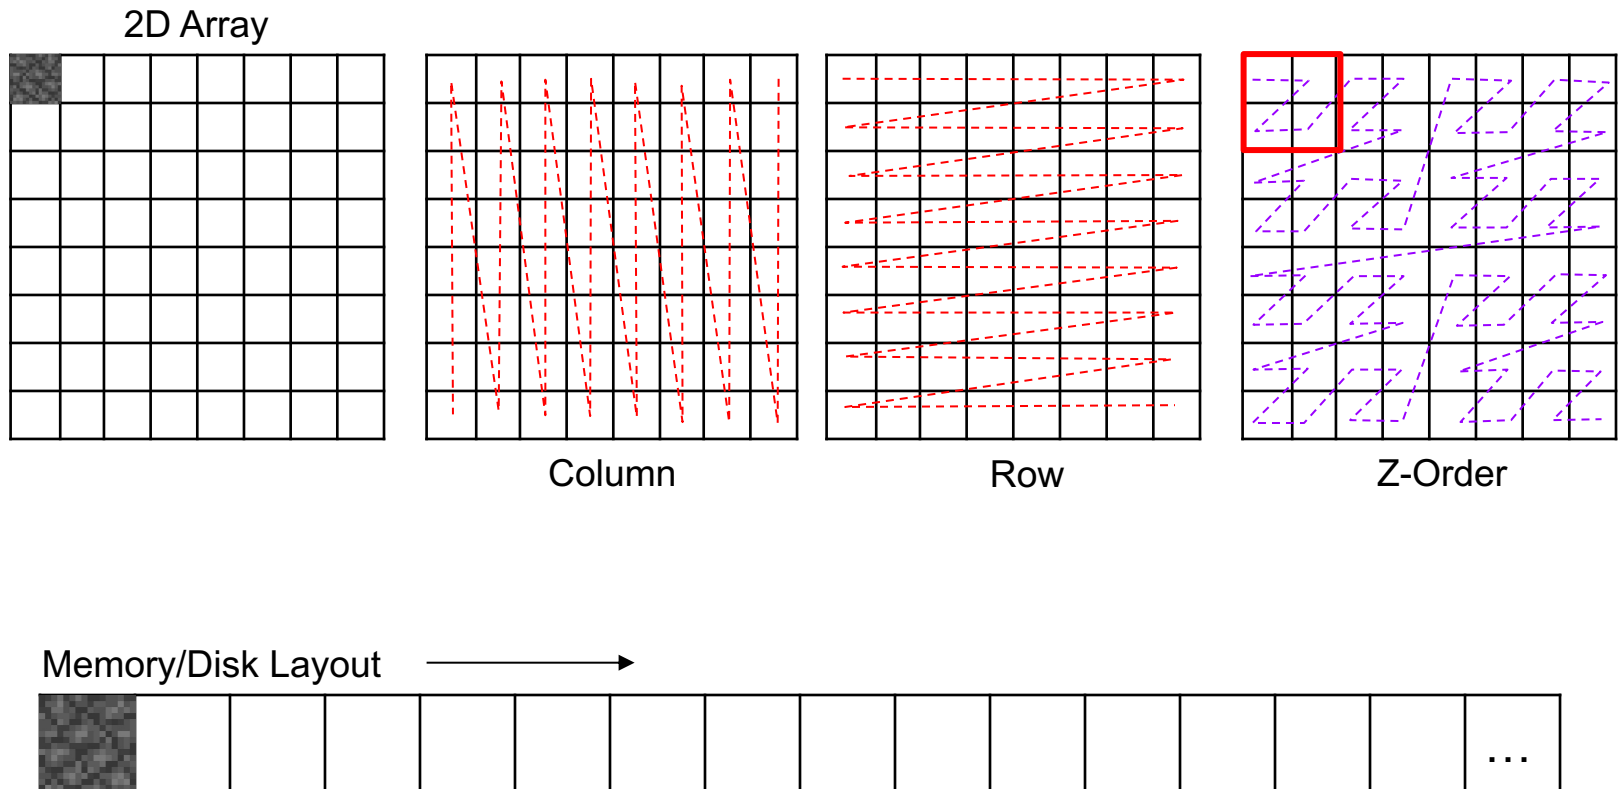
- Graph DBMS
  - Introduction
  - Storage and Indexing
  - Partitioning
- Array DBMS
  - Introduction
  - Physical Array Storage
  - Space filling curves

# Multi-dimensional Arrays

---



- Scientific data is often in the form of multi-dimensional arrays (matrices).
  - Arrays of images, video frames, or various scientific data
- Typical workload
  - Mostly analytical: slice, dice, aggregate
  - Algorithms: anomaly detection, object detection, etc.
- Industry data file format: HDF5 (Hierarchical Data Format)

# Physical organization of an Array



# Space Filling Curves

- A space-filling curve is a mathematical curve that passes through every point in a 2D (or higher) dimensional space.
- Maps 1D line into a high-dimensional space (and vice versa)
  - Using mathematical formulae
- Examples: Hilbert-curve (images below\*) / Z-Order / Peano-curve



\*<https://github.com/jakubcerveny/gilbert>

# Agenda

---

- Graph DBMS
  - Introduction
  - Storage and Indexing
  - Partitioning
- Array DBMS
  - Introduction
  - Physical Array Storage
  - Space filling curves
- Q&A
- Midterm this Thursday during class time.