

Big Data Systems

Djellel Difallah

Spring 2023

Lecture 8 – Distributed File Systems

GFS and HDFS

So far ...

- We looked at **Data Management** techniques
 - RDBMSs
 - Relational Schema
 - NoSQL
 - Mostly simple operations such as put/get
- Now we look at very large scale **Data Processing** techniques
 - Massive amounts of data
 - Read mostly workloads (OLAP)
 - Efficiency (data processed per second) over consistency
 - **Separating storage and computation**

Outline

- Introduction
- The Google File System
 - Motivation
 - Design Principles
 - Architecture and Operations
- The Hadoop Distributed File System
 - Architecture
 - Features

Introduction

- Your job is to *process* 1TB of data
 - You: *Sure!*
 - But first, you need to read it.

```
with open(bigFile) as f:  
    line = f.readlines()  
    ...
```
 - You: ☹ It's taking forever.
- HDD (~ 150 MB/s) ~ 2.5 hours
- SSD (~ 500 MB/s) ~ 1 hour
- NVMe: couple of minutes.
- Multithreading will not help (*especially not for HDD*)

Introduction

- That was an exercise:
 - 10TB, 100TB, 1PB ??
- The data storage medium is clearly a bottleneck
- We need a way to increase the throughput
 - Yes: use multiple disks
 - Not scalable on a single machine (Bus limitations)
- N Machines with K disks
 - Cost effective when using cheap machines.
 - Hard to program


File System

- **Metadata** = information about the file
 - Includes name, access permissions, timestamps, file size, & locations of data blocks
 - **inodes** (index **nodes**) are data structures containing meta information about the file
- **Data** = actual file contents
- A file system stores all the **bits and pieces** of a given file on the same device.
- For example: Linux directories are just files that contain lists of names and inodes


RAID: Redundant Array of Independent Disks

- RAID 0 (Striping):
 - Data is **striped** across multiple disks, which improves read and write performance.
 - RAID 0 does not provide any redundancy or fault tolerance: data loss will occur if a single disk fails.
- RAID 1 (Mirroring):
 - Data is **duplicated** across two or more disks, providing fault tolerance through redundancy.
 - RAID 1 improves read performance but has a higher cost due to the need for additional storage.
- RAID 5 (Striping with Parity):
 - Data and **parity information** are striped across three or more disks.
 - Parity is an error detection and correction techniques.
 - RAID 5 provides fault tolerance and improves read performance but requires extra storage for parity data.
 - A minimum of 3 disks, and can handle 1 disk failure.

RAID Examples



RAID Controller



Client Server File Systems

- Network Attached Storage (NAS)
 - Shared Disk Architecture: Single point of failure.
 - Point of congestion
 - Industry solution:
 - Replication, recovery and client caching



Parallel File Systems

- Distributed Architecture
 - **Multiple servers** work together to store, manage, and access data
 - Enhance performance
- Data Chunking (Striping)
 - Data is divided into smaller **chunks** (stripes) and distributed across multiple storage devices
 - **File data can be on multiple servers**
 - Enables concurrent reading and writing for improved performance
- Metadata Management
 - **Metadata can be stored and managed separately from data (stored on a different node)**
 - Allows for efficient searching, indexing, and data organization

Parallel File Systems (cont.)

- Scalability
 - Designed to handle large-scale storage and computing environments
 - Can easily expand storage capacity and performance as needed
- Fault Tolerance & Data Redundancy
 - Built-in mechanisms to protect against data loss and hardware failures
 - Examples: replication, erasure coding, and RAID techniques
- Load Balancing
 - Ensures optimal distribution of data and workload across available resources
 - Minimizes bottlenecks and maximizes performance

Consideration

- Component (Disk) failure is the norm
 - File system = thousands of storage machines
 - Some % not working at any given time
- Most files are large (not small).
 - Block size and I/O operations
- The programming models changes
 - Rethink existing solutions

Programming Model

- **Shared Memory**
 - Threads
 - Shared data structures with locks etc.
 - Easy to reason about.
 - The memory becomes a limitation
- **Message Passing** (many techniques here)
 - Master-slave
 - Producer-consumer
 - Very difficult to program correctly and to reason about.
 - Every new task will require reengineering

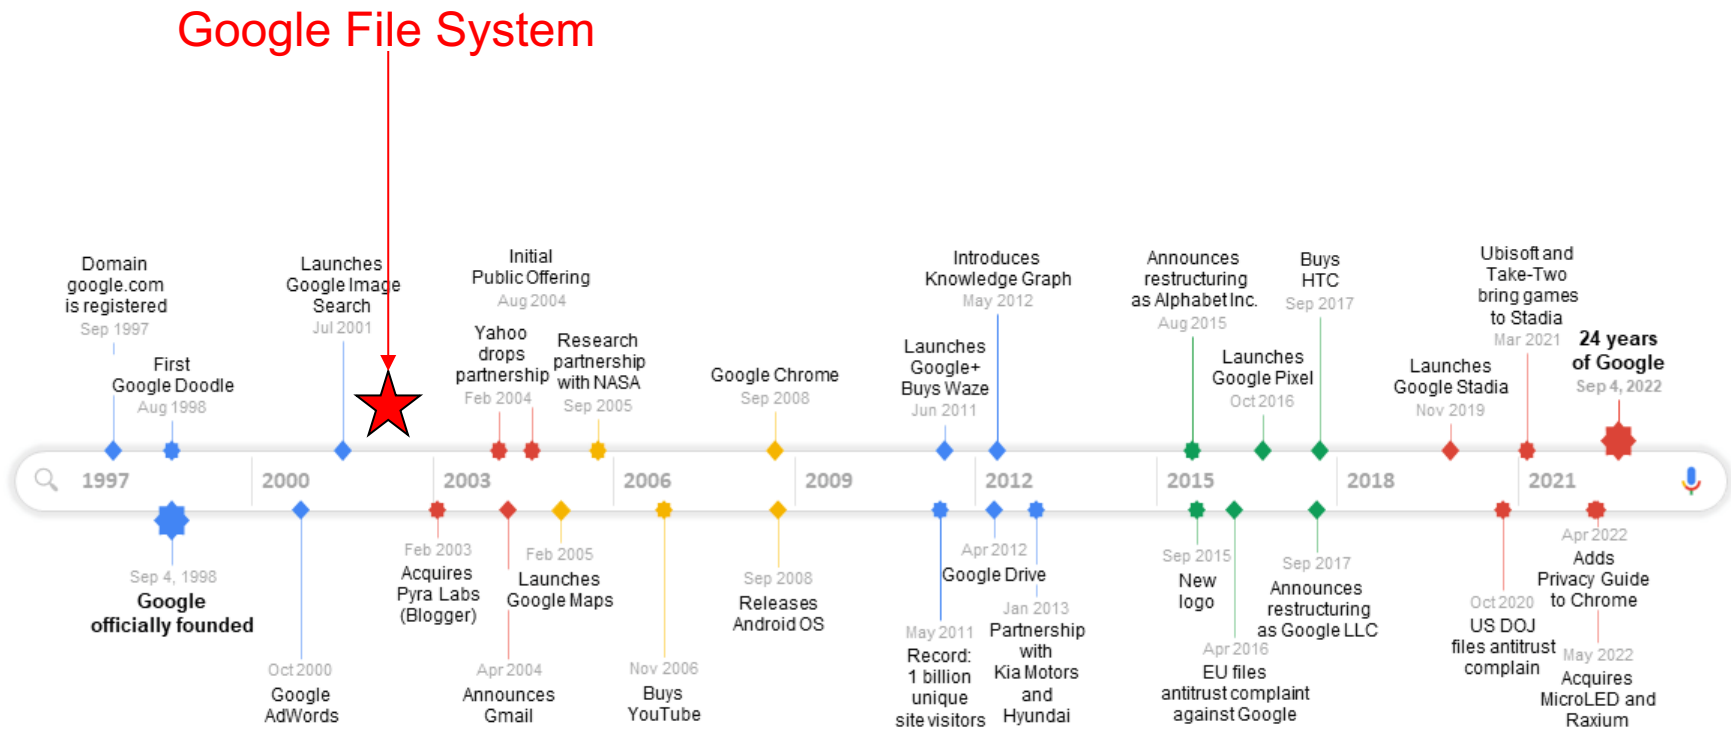
The Google File System

GFS

Introduction

- New parallel file system to meet the needs of Google's large-scale data processing challenges
- *"The Google File System" (2003) by Ghemawat et al.*
 - The paper describes the design and implementation of the Google File System a scalable, distributed file system.

Google History Context



Why are we learning about Google File System?

- Covers important (and timeless) themes in distributed systems:
 - parallel performance, fault tolerance, replication, consistency
- Practical: a successful real-world implementation and use case.
 - [MapReduce](#), BigTable built on top of GFS
- Large scale
- Demonstrate the use of [weak consistency](#)
- [Single master](#) architecture
- Simple but thorough system paper to read

Motivation

- Commodity Hardware
 - Error is frequent in large setups
- Large Files
 - Multiple 100Mb+ files
- File operations
 - Read and Append
 - Mostly sequential reads

Design Principles

- Fault tolerance:
 - GFS is designed to handle faults gracefully, including hardware failures and network issues.
 - It achieves this through replication, checksums, and automatic recovery mechanisms.
- Data replication:
 - To ensure durability and availability, the system replicates each chunk across multiple node (chunk servers), default is three replicas.
- High throughput:
 - GFS is optimized for **high sustained throughput** over **low latency**,
workloads typically involve large data processing tasks.

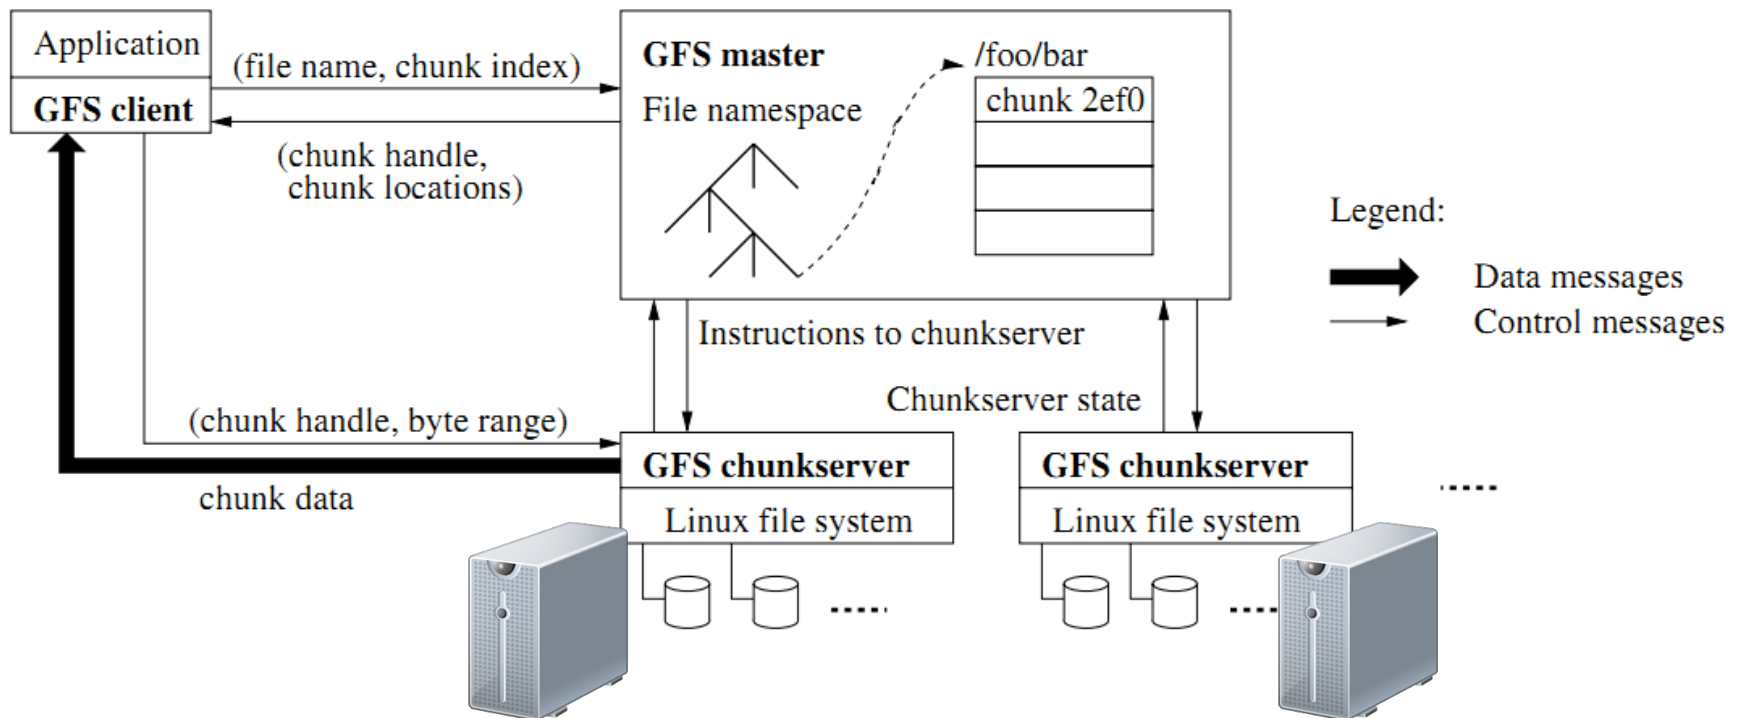
Design Principles

- Consistency and atomicity:
 - GFS provides a consistent view of the file system by using a combination of techniques like write serializability, chunk versioning, and atomic record appends.
- Chunk-based storage:
 - GFS stores files as fixed-size chunks and uses a single master server to manage the metadata associated with these chunks.
 - The actual chunk data is stored on chunk servers distributed across the cluster.

Design Principles

- Scalability:
 - The architecture is designed to support massive scale, allowing it to handle large numbers of files and a vast amount of storage.
 - Runs on commodity hardware
- Master-Slave Architecture
 - A single node (master server) is responsible for managing the file system metadata, such as the namespace, access control information, and chunk location information.
 - The master server also handles garbage collection and chunk migration to balance load and recover from failures
 - Single point of failure

GFS Architecture



Master

- Stores the full file system metadata
 - Namespace (i.e., file system tree)
 - Access control info
 - Filename to chunks mappings
 - Current locations of chunks
- Manages
 - Chunk leases (locks)
 - Garbage collection (freeing unused chunks)
 - Chunk migration (copying/moving chunks)
- Fault tolerance
 - **Operation log** replicated on multiple machines
 - New master can be started if the master fails by replaying the operation log.
- Periodically communicates with all chunkservers – Via heartbeat messages to get state and send commands
 - HeartBeat is a short message:
 - Ensure that server is alive
 - Collect statistics
- Advantage
 - Given the setup (read mostly) this architectural choice simplifies the system immensely.

File Chunks and Chunkserver

- Chunk size
 - 64 MB (default)
 - DBMS to Big Data: From page size to chunk size
- Chunk Handle:
 - A global unique identifier assigned by the master
- Chunkservers
 - Stores chunks on local disks as Linux files
 - Stores a 32-bit checksum with each chunk in memory and logged to disk (detect data corruption)

Metadata

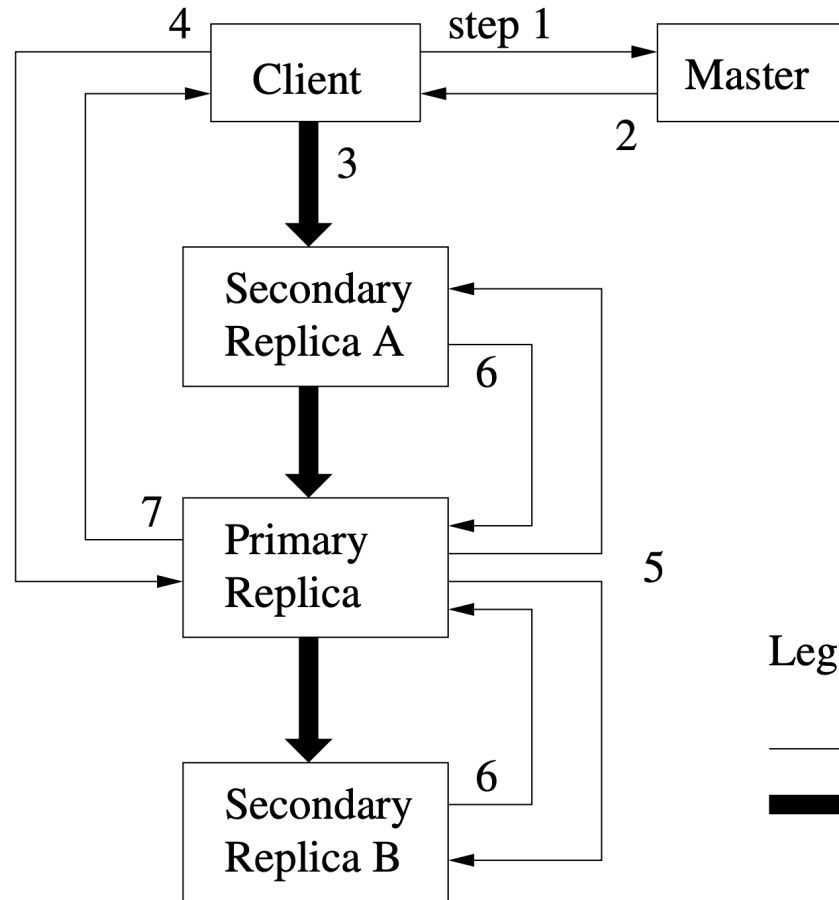
- All metadata is stored in the **master system memory** for fast operations and efficient background scans
 - Background scans are processes used for maintenance, e.g., chunk garbage collection.
- Types of metadata:
 - Namespace: Hierarchical information about files and directories
 - Filename: an array of chunk handlers
 - Chunk handler: list of chunk servers, primary chunk server, chunk version number
- First 2 types of metadata are persisted by logging mutations to operation log (oplog) on the master's local disk
 - Operation log replicated periodically to a remote machine (shadow master)
- Master does not store chunk location information **persistently**
 - Chunkservers provide chunk location and primary status information on master startup or when joining the cluster

Consistency

- Atomic operations:
 - Namespace **mutations** are atomic and managed by the master
 - Namespace locking ensures atomicity
 - Master's operation log defines a global order
 - Record append operation is atomic
- Chunk replication:
 - Ensures data consistency after mutations
 - Applies mutations in order by designating a primary replica
 - Uses chunk version numbers to detect stale replicas
 - Managed by the master
 - If stale replicas are not used for mutations or provided to clients; garbage collected when possible

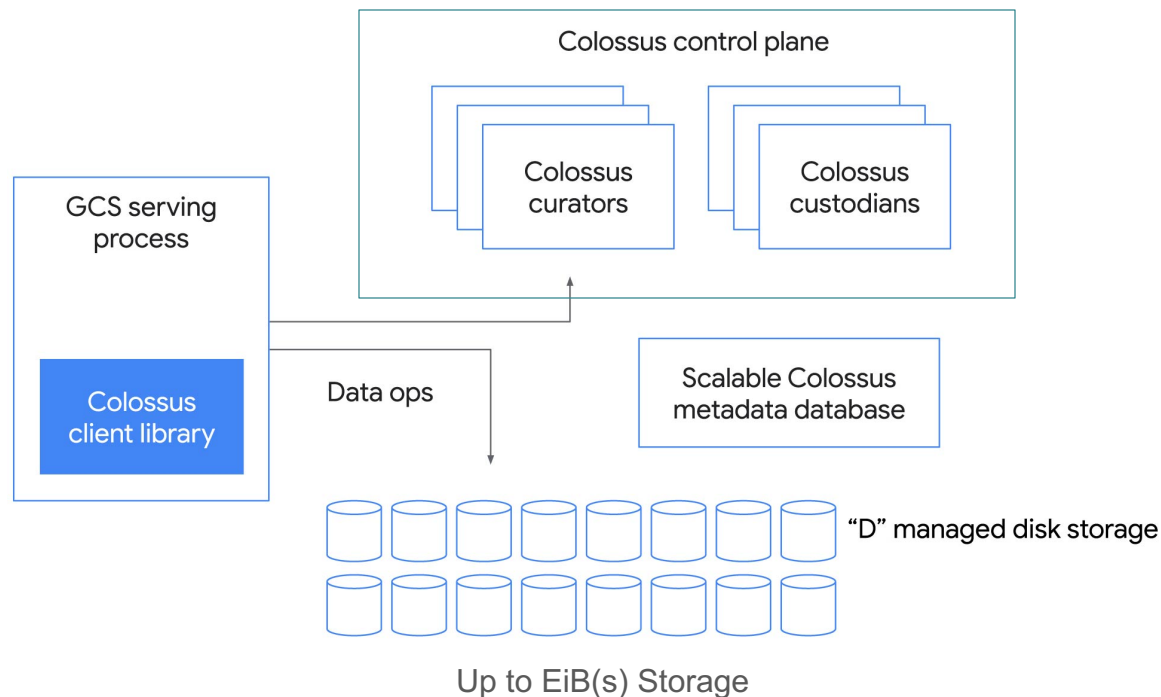
Write Data Flow

1. Lease acquisition
 - Chunk Version Increment
2. Primary selection
3. Data Forwarding
4. Data buffering and Ack
5. Write Command
 - Primary Mutation
 - Secondary Mutation
6. Primary Ack
7. Client Ack



Google next generation GFS: Colossus

- Google Colossus: Better scalability and availability
 - Blog post: ["A peek behind colossus googles file system"](#)
 - Curator: Distributed metadata (NOSQL)
 - Custodian: storage management (space balancing, raid, etc.)





Apache Hadoop Distributed File System

HDFS

Apache Hadoop

- A Platform (Ecosystem) for Big Data processing
 - Started by Yahoo! and was inspired by the following papers from Google
 - MapReduce

Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: simplified data processing on large clusters." Communications of the ACM 51.1 (2008): 107-113.
 - GFS

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." ACM SIGOPS operating systems review. Vol. 37. No. 5. ACM, 2003.
- Written *mostly* in Java
- Main core components
 - Hadoop Distributed File System (HDFS)
 - A Map Reduce implementation
 - Other management services such as *YARN*

HDFS Motivation

- Lots of commodity hardware
 - Cheap PCs
- High failure rates
 - Commodity components failure is much higher than server grade.
- Write-once read many times data
 - Logs, archives etc.
 - Processing is mostly read oriented
- Large streaming reads over random access
 - High sustained throughput over low latency

HDFS Features

- A stand-alone general purpose Distributed File System
 - used as a part of a Hadoop cluster
- Highly scalable
 - Many production clusters with Petabytes of data
- Fault tolerant
- Offers a simplified data access
 - The users see their file (data) as if they were on a single machine
 - default configuration is well suited for many installations
 - Tuning is required for large deployments
- Linux like security, **limited to file level permission**

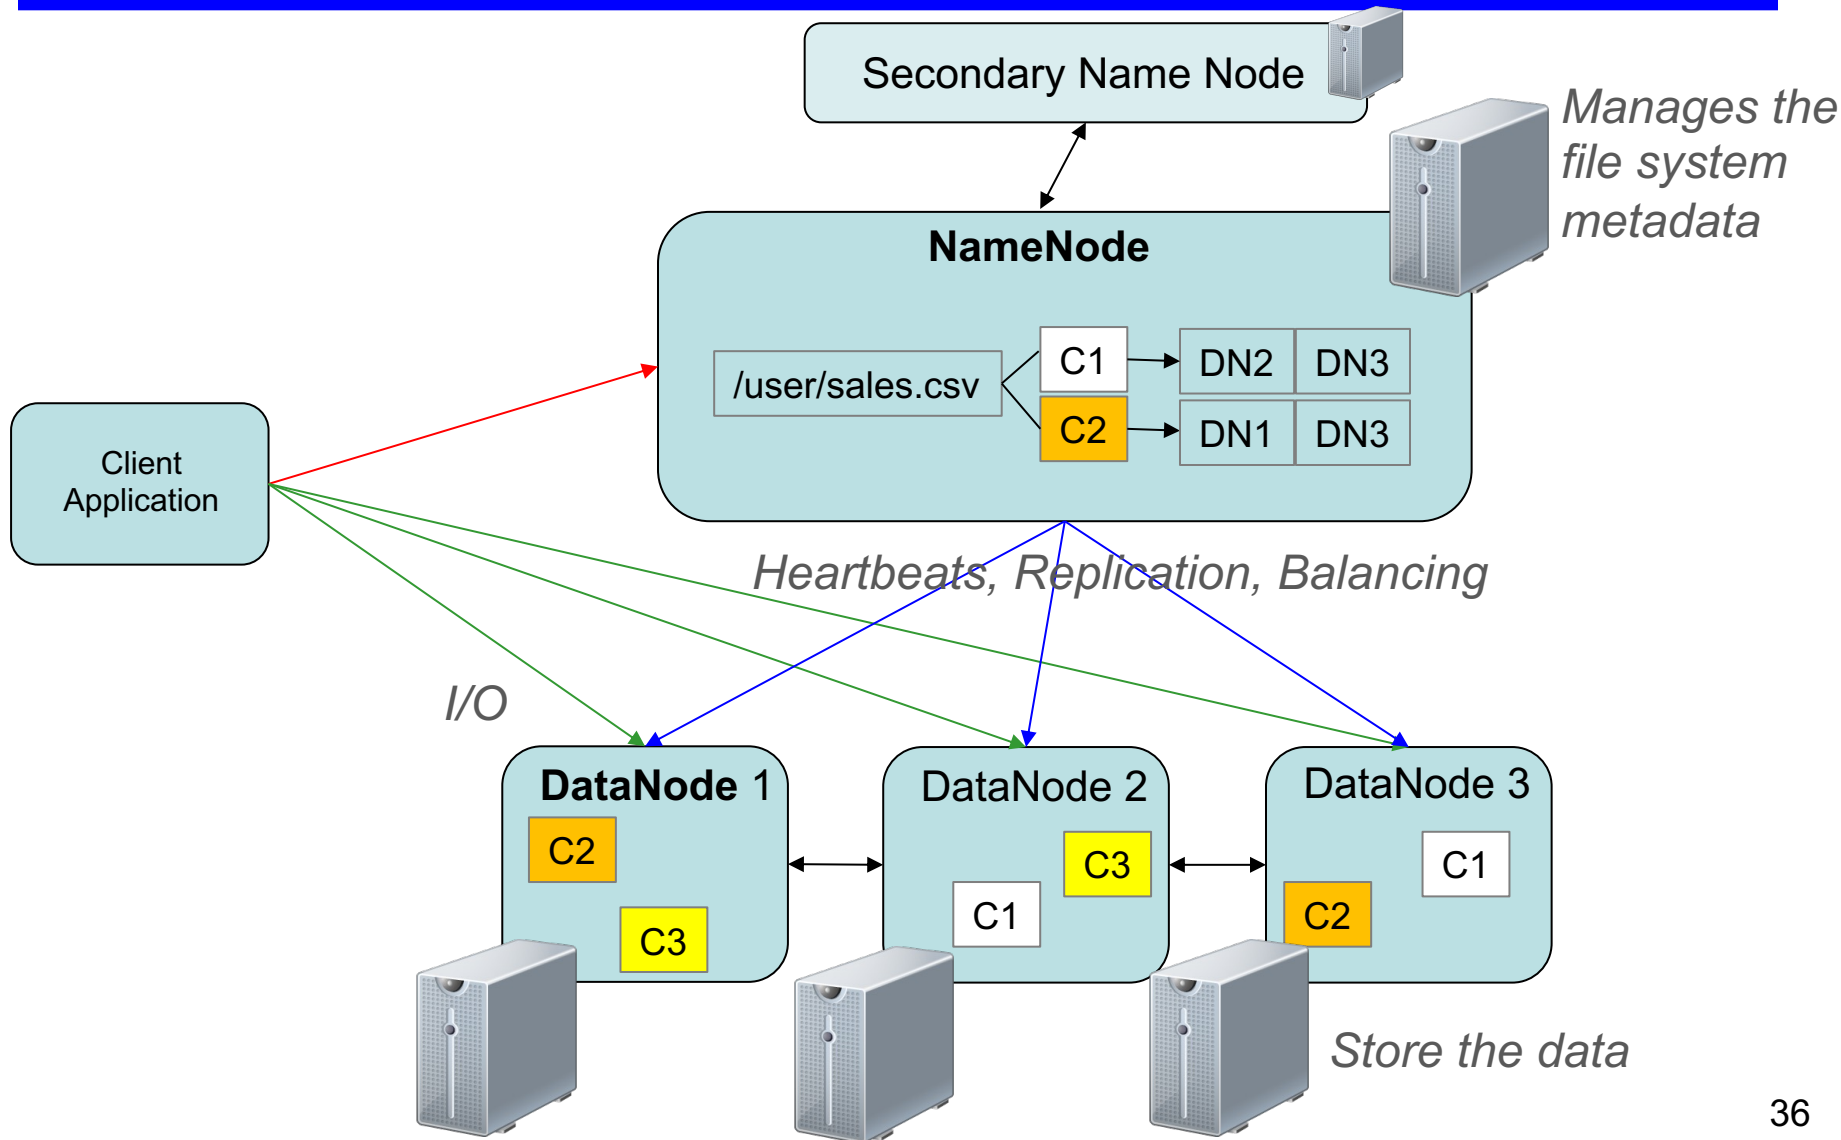
HDFS Basic Operations

- When an input file is added to HDFS
 - File is split into smaller **blocks** of **fixed size**
 - Each block is replicated to multiple hosts (machines)
 - Each replicated block is stored on a different host
 - Replication level is configurable
 - Default is 3
 - If a host crashes or is decommissioned
 - All blocks are replicated to a new host
 - In case a new host is added
 - Blocks will be rebalanced (avoid data skew)

Blocks

- Single unit of storage
- Transparent to the end-use
- Block is traditionally 64/128/256 MB
- Block size is fixed whether you store a file of 64MB or less.
- It is large enough to force a tradeoff between a seek and sequential disk read.
 - Time to read a block = seek time + transfer time
 - We want to minimize $\frac{\text{Seek Time}}{\text{Transfer Time}}$

HDFS Architecture



Name Node

- Managing the file system namespace:
 - Holds file/directory structure, metadata, file-to-block mapping, access permissions, etc.
- Coordinating file operations:
 - Directs clients to datanodes for reads and writes
 - No data is moved through the namenode
- Maintaining overall health:
 - Periodic communication with the datanodes
 - Block re-replication and rebalancing
 - Garbage collection
- Executes file system namespace operations like opening, closing, and renaming files and directories

Data Node


- Actual storage and management of data block on a single host
- Provides clients with access to data
- Perform block creation, deletion, *and* replication upon instruction from the NameNode
- Sends heartbeat (every 3 seconds) and a blockreport (6hours) to NameNode

Secondary Name Node

- The name node keep the current state of the HDFS in an image file (`fsimage`) which it reads during the startup
- All modifications are written to a log (`edits`)
- NN merges `fsimage` and `edits` during startup
- The Secondary NameNode is not a failover NN
 - Role: check pointing
 - It merges periodically `fsimage` and `edits`
 - Every 1 Hour or 1 Million transactions by default (configurable)

Replication

- Maximize the reliability, availability and bandwidth
 - Opportunistic execution
- Replicas spread across machines and racks
- Consider geo replication for disaster recovery
- The master determines replica placement. In HDFS:
 - 1st replica on local node (w.r.t writing process) or random node
 - One replicas is usually placed on the same rack as the node writing to the file cross-rack network I/O is reduced
 - One replica on a different node in a different rack



Rebalancer

- HDFS
 - May produce non-uniform distribution of blocks across hosts
 - e.g., new machines added or removed
 - Does not consider different hardware generations
 - Does not consider current hardware utilization
 - Operates on block level not on file level
- HDFS architecture is compatible with data rebalancing schemes
- The system administrator can perform periodic cluster wide rebalancing of the blocks.

`hadoop balancer <threshold>`

- The threshold (X%) sets a target for the balance state (default 10)
- **Each** node's utilization ratio is within X% of the overall cluster utilization
 - Utilization Ratio = $\frac{utilizedSpace}{TotalCapacity}$
- Smaller threshold leads to more balanced distribution
 - Might take a lot of time, esp. when concurrent utilization

Block Placement Policy

- HDFS allows the implementation of custom block placement policies
- Alternative Placement Policies, e.g.,:
 - Use lower utilization servers first
 - Operate on file level
 - Assign weight to DataNodes
 - Move more blocks to newer generation hardware

