

# Big Data Systems

Djellel Difallah

Spring 2023

Lecture 10 – Apache Spark

# Outline

---

- Introduction
- Apache Spark Motivation
- RDD: Resilient Distributed Datasets
- Spark Internals

# MapReduce Problems

---

- Many problems aren't easily mapped into a MapReduce job
- Persistence to disk is typically slower than in-memory processing
  - Shuffle phase is disk intensive
- Jobs reload data from disk storage on each new execution
  - No-reutilization

# Motivation

---

- Towards “distributed data programming”
  - **Iterative** data processing
    - Multiple runs of a Map/Reduce program
  - **Interactive** data processing with intermediary data reuse
    - Arbitrary code + parallel data processing
- Many **specialized** frameworks on top of M/R have been created
- **Increasingly better** hardware in Hadoop deployments:
  - High-speed networks
  - Larger memory capacity



# APACHE SPARK

# Design Ideas

---

- Retain the attractive properties of MapReduce
  - Scalability
  - Data locality
  - Fault tolerance
- Support more operation
- **Lesson learned from other systems:** Execution are DAGs (Directed A-cyclic Graphs) of tasks.
  - Unification has benefits for user (learning curve) and the system (code base, complexity etc.)

💡 Keep intermediary/computed data **in-memory**

# Apache Spark

---

A general purpose data processing engine

- Defines a large set of operations (as opposed to simple “map” and “reduce”)
- Operations can be arbitrarily combined in any order
- Programming at a higher level of abstraction; work with distributed dataset as if it was local
- Combines multiple data processing types (SQL, ML, Graph)

# Getting Started with Spark

---

- Install dependencies:
  - Java and Scala
- On Mac:
  - brew install apache-spark
- On Windows
  - [Tutorial](#)



# Interact with Spark using Scala or PySpark shell

```
dd115@ADUAEI12744LPMX ~ % spark-shell
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/04/11 10:04:25 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Spark context Web UI available at http://10.225.94.244:4040
Spark context available as 'sc' (master = local[*], app id = local-1681193065734).
Spark session available as 'spark'.
Welcome to

  ____      _
 / ___|  _ \| | | |
 \___ \| |_) | |_| |
  ___) | |_) | | | |
 |____|_|_|\___|_|_|_|
version 3.3.2

Using Scala version 2.12.15 (OpenJDK 64-Bit Server VM, Java 20)
Type in expressions to have them evaluated.
Type :help for more information.

scala> 23/04/11 10:04:37 WARN GarbageCollectionMetrics: To enable non-built-in garbage collector(s) List(G1 Concurrent GC), users should configure
it(them) to spark.eventLog.gcMetrics.youngGenerationGarbageCollectors or spark.eventLog.gcMetrics.oldGenerationGarbageCollectors

[scala>
[scala>
scala> █
```

```
dd115 — java • Python — 146x28

Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/04/11 10:30:06 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Welcome to

  ____      _
 / ___|  _ \| | | |
 \___ \| |_) | |_| |
  ___) | |_) | | | |
 |____|_|_|\___|_|_|_|
version 3.3.2

Using Python version 3.11.3 (main, Apr 7 2023 19:29:16)
Spark context Web UI available at http://10.225.94.244:4040
Spark context available as 'sc' (master = local[*], app id = local-1681194607212).
SparkSession available as 'spark'.
>>> █
```



Zaharia, Matei, et al. "Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing." In Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation (NSDI'12)

# RESILIENT DISTRIBUTED DATASET

# Apache Spark RDD

---

- Resilient Distributed Dataset (RDD)

*“Resilient Distributed Datasets (RDDs) are a distributed memory abstraction that lets programmers perform in-memory computations on large clusters in a fault-tolerant manner.” (Zaharia 2012)*

- In other words: A Distributed Data Collection

- Is this sharding? 🤔
- mylist = [1, 2, 3]
- mylistRDD = [1, 2, 3]

# RDD Characteristics

---

- **In-Memory first**
- **Immutable or Read-Only**
- **Lazy evaluated**
- **Cacheable**
- **Parallel**
- **Typed**
- **Partitioned**

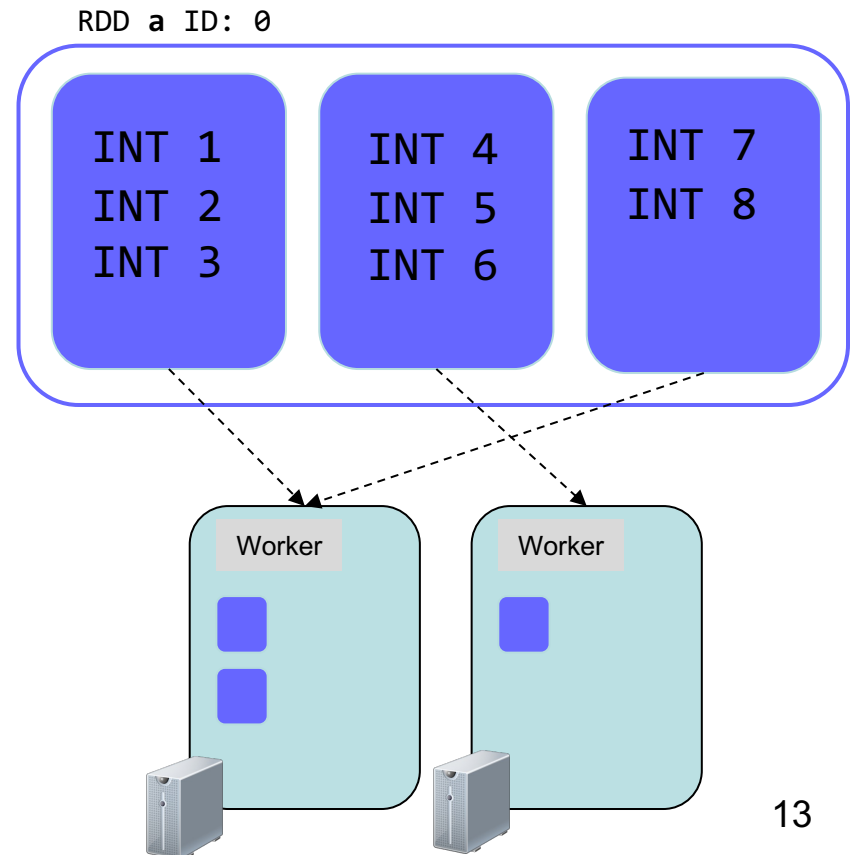
# RDD Creation and Partition

An RDD can be created in 2 ways:

- Parallelize a collection (list, set, dictionary, etc.)
- Read data from an external source (HDFS, S3, etc.)

```
var firstRDD = sc.parallelize(1 to 8)
firstRDD.cache()
firstRDD.count()
```

```
firstRDD = sc.parallelize(range(1, 9))
firstRDD.cache()
firstRDD.count()
```



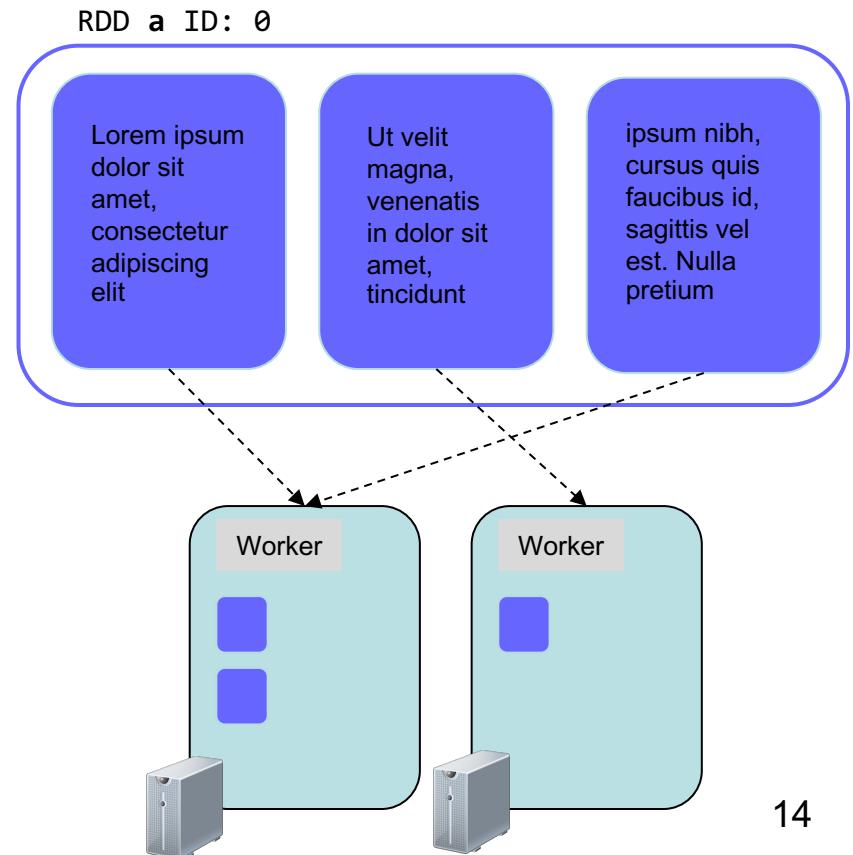
# RDD Creation and Partition

An RDD can be created in 2 ways:

- Parallelize a collection
- Read data from an external source (HDFS, S3, etc.)

```
var secondRDD = sc.textFile("input")  
secondRDD.cache()
```

```
secondRDD = sc.textFile("input")  
secondRDD.cache()
```



# RDD Operation and Lifecycle

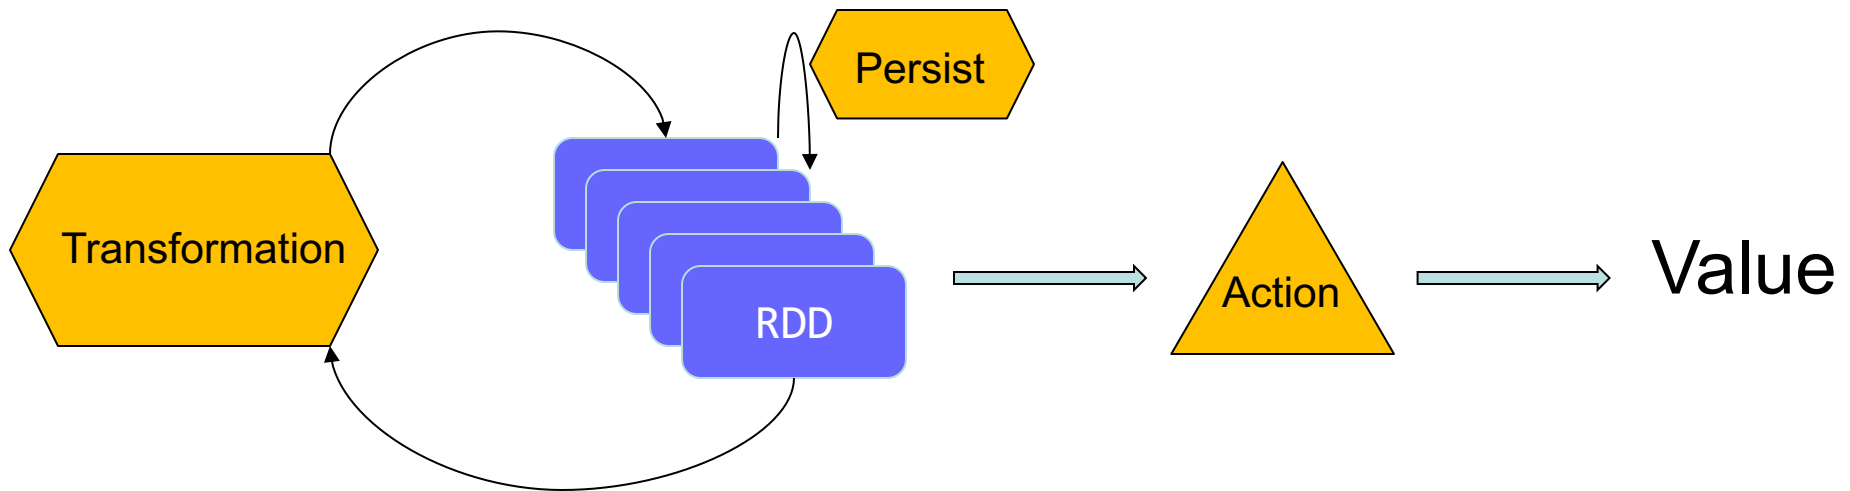
---

- Transformations

- Lazy operations that return another RDD

- Actions

- Operations that trigger computation and return values



# Example in PySpark

---

Find the number of distinct *names* by “first letter”

```
# Example
# input : alba, david, boyl, doris, bob, brave
# output: (d,2), (b,3), (a,1)

input = sc.textFile("hdfs://names")
tuple = input.map(lambda name: (name[0], name))
values = tuple.groupByKey()
counts = values.mapValues(lambda name: len(set(name)))
counts.collect() # Action!
```



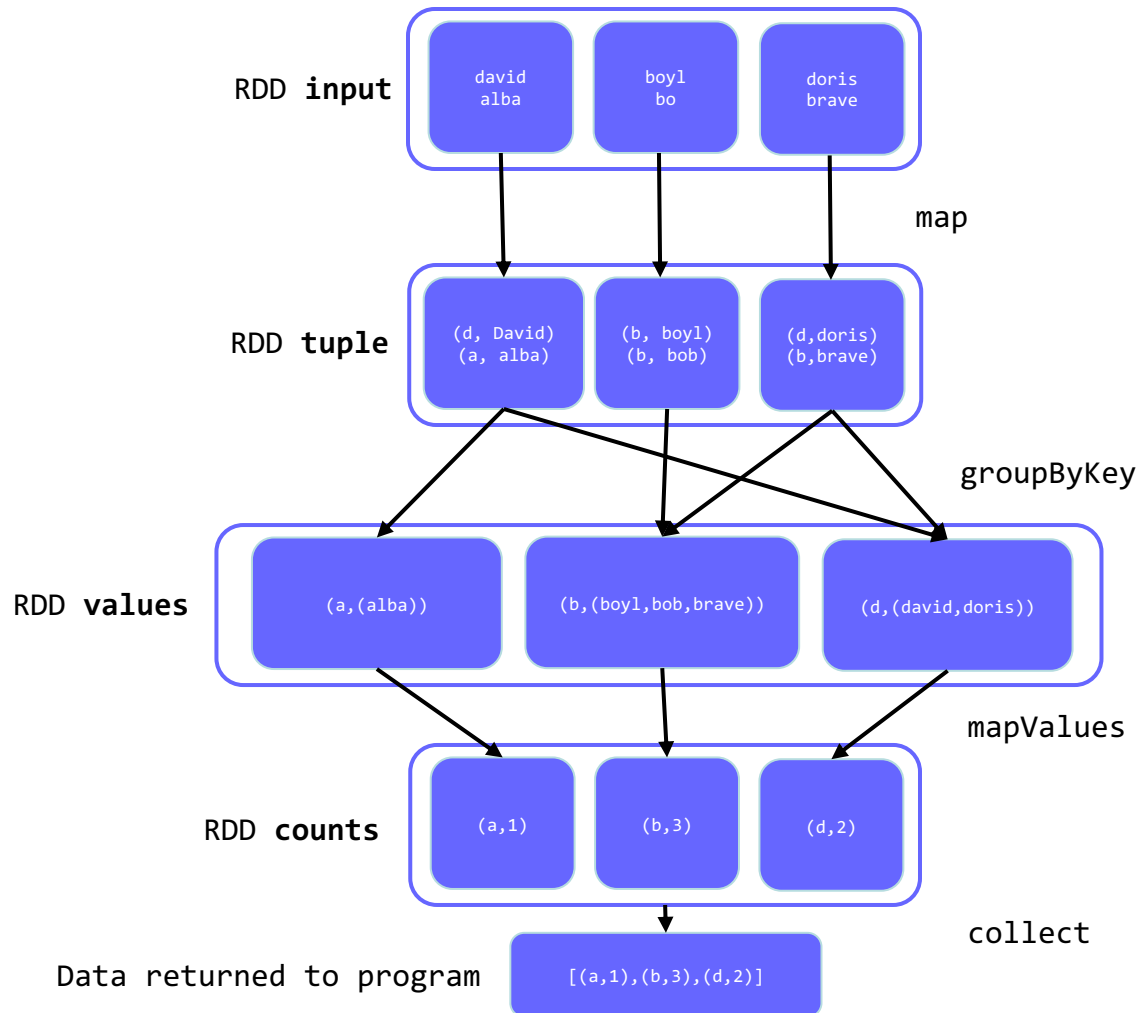
# Example in Scala

---

Find the number of distinct *names* by “first letter”

```
// Example
// input : alba, david, boyl, doris, bob, brave
// output: (d,2), (b,3), (a,1)
var input = sc.textFile("hdfs://names")
var tuple = input.map(name => (name.charAt(0), name))
var values = tuple.groupByKey()
var counts = values.mapValues(name => name.toSet.size)
counts.collect() // Action !
```

# Example RDD



# Misc. Examples

---

```
###
x = sc.parallelize([1,2,3])
y = x.map(lambda x: (x, x**2))
x: [1, 2, 3]
y: [(1, 1), (2, 4), (3, 9)]

###
x = sc.parallelize([('B',5),('B',4),('A',3),('A',2),('A',1)])
y = x.groupByKey()
x: [('B', 5), ('B', 4), ('A', 3), ('A', 2), ('A', 1)]
y: [('A', [3, 2, 1]), ('B', [5, 4])]

###
x = sc.parallelize([('A',(1,2,3)),('B',(4,5))])
y = x.mapValues(lambda x: [i**2 for i in x])
x: [('A', (1, 2, 3)), ('B', (4, 5))]
y: [('A', [1, 4, 9]), ('B', [16, 25])]

###
x = sc.parallelize([1,2,3])
y = x.collect()
y: [1, 2, 3]
```

# RDD Operations

---

## Transformations

(define a new RDD)

map()  
flatMap()  
distinct()  
filter()  
groupByKey()  
reduceByKey()  
coalesce()  
sortByKey()  
partitionBy()  
sample()  
join()  
union()  
...

**persist()**

**cache()**

*Special transformations: Mark the RDD for persistence*

## Actions

(return results to program)

reduce()  
collect()  
saveAsTextFile()  
count()  
first()  
take(n)  
countByKey()  
takeSample()  
foreach()  
...

# RDD Types

---

HadoopRDD

JdbcRDD

JsonRDD

SchemaRDD

ShuffledRDD

UnionRDD

CassandraRDD

...

Specialized RDDs; Check out the code repository:

<https://github.com/apache/spark/tree/master/core/src/main/scala/org/apache/spark/rdd>

# RDD Interface

---

- Set of partitions (“splits”)
- List of dependencies on parent RDDs
- Function to compute a partition given parents
- **Optional** preferred locations
- **Optional** partitioning information for Key/Value RDDs (Partitioner)

Base RDD code:

<https://github.com/apache/spark/blob/master/core/src/main/scala/org/apache/spark/rdd/RDD.scala>

# Example: HadoopRDD

---

- `partitions` = one per HDFS block
- `dependencies` = none
- `compute(partition)` = read corresponding block
- `preferredLocations(part)` = HDFS block location
- `partitioner` = none

# RDD Manipulation

---

Users can control two aspects of RDDs:

- *Persistence* (in RAM, reuse)
- *Partitioning* (*hash*, *range*, [ $<k, v>$ ])

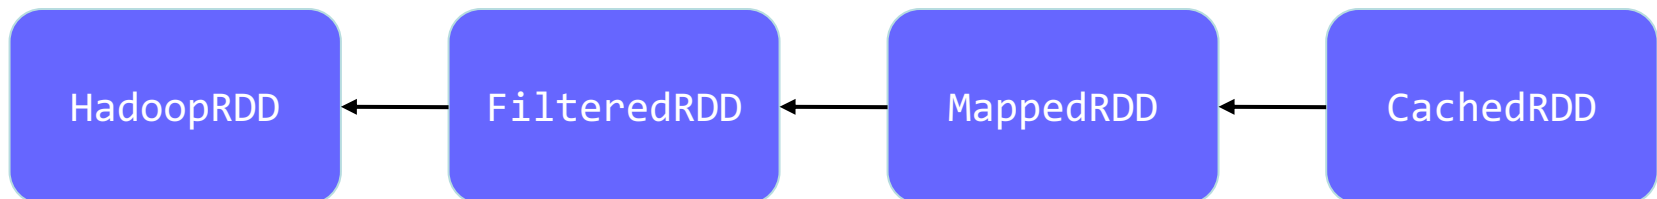


# RDD Lineage

---

- **Lineage:** the sequence of RDDs (Resilient Distributed Datasets) that form the dependencies between the RDDs in a Spark application
- **Fault Tolerance:** Upon node failure RDDs recompute lost data by reapplying the transformations used to build them

```
var errors = sc.textFile("hdfs://logs")  
  .filter(_.contains("error"))  
  .map(_.split('\t')(2))  
  .cache()
```



# RDD vs. Distributed Shared Memory

---

Concern	RDDs	Distr. Shared Mem.
Reads	Fine-grained	Fine-grained
Writes	Bulk transformations	Fine-grained
Consistency	Trivial (immutable)	Up to app / runtime
Fault recovery	Fine-grained and low-overhead using lineage	Requires checkpoints and program rollback
Straggler mitigation	Possible using speculative execution	Difficult
Work placement	Automatic based on data locality	Up to app (but runtime aims for transparency)

# Benefits of RDD Model

---

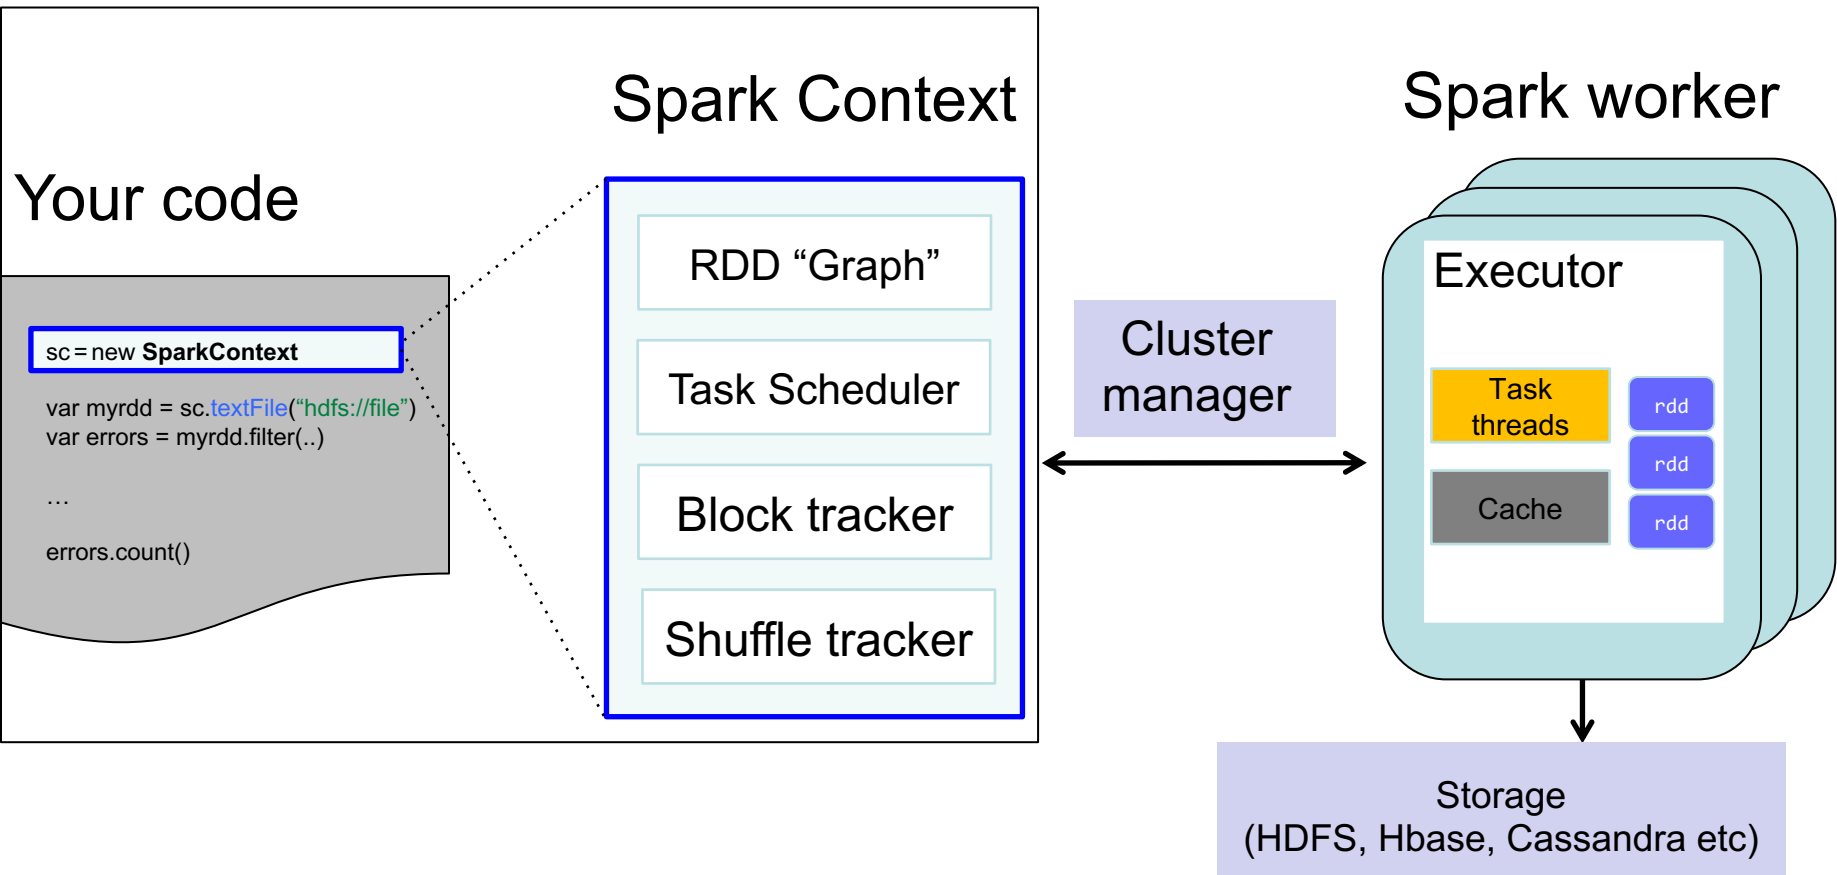
- Consistency is easy due to immutability
- Inexpensive fault tolerance (log lineage rather than replicating/checkpointing data)
- Locality-aware scheduling of tasks on partitions
- Despite being restricted, model seems applicable to a broad variety of applications



# APACHE SPARK INTERNALS

# Spark Components

## Driver



# Spark Execution Modes

---

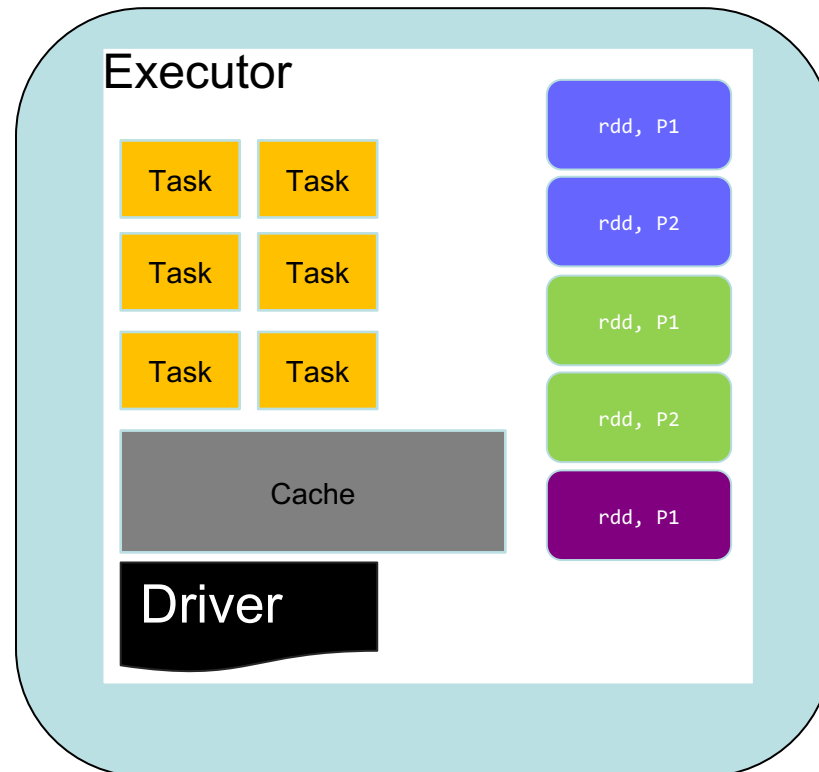
- Local (Local machine)
- Standalone “Cluster) (manually configured cluster)
- YARN (Hadoop cluster)

Using Container Orchestration Engines:

- Mesos
- Kubernetes

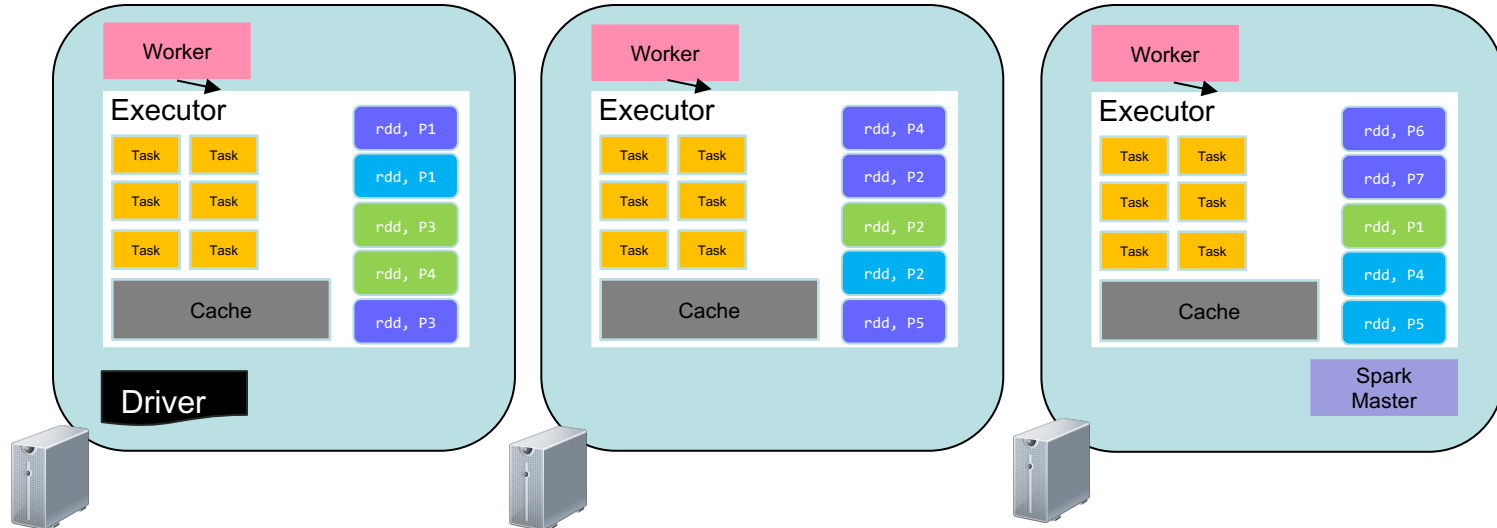
# Local Mode

---



`./bin/spark-shell --master local[6]`

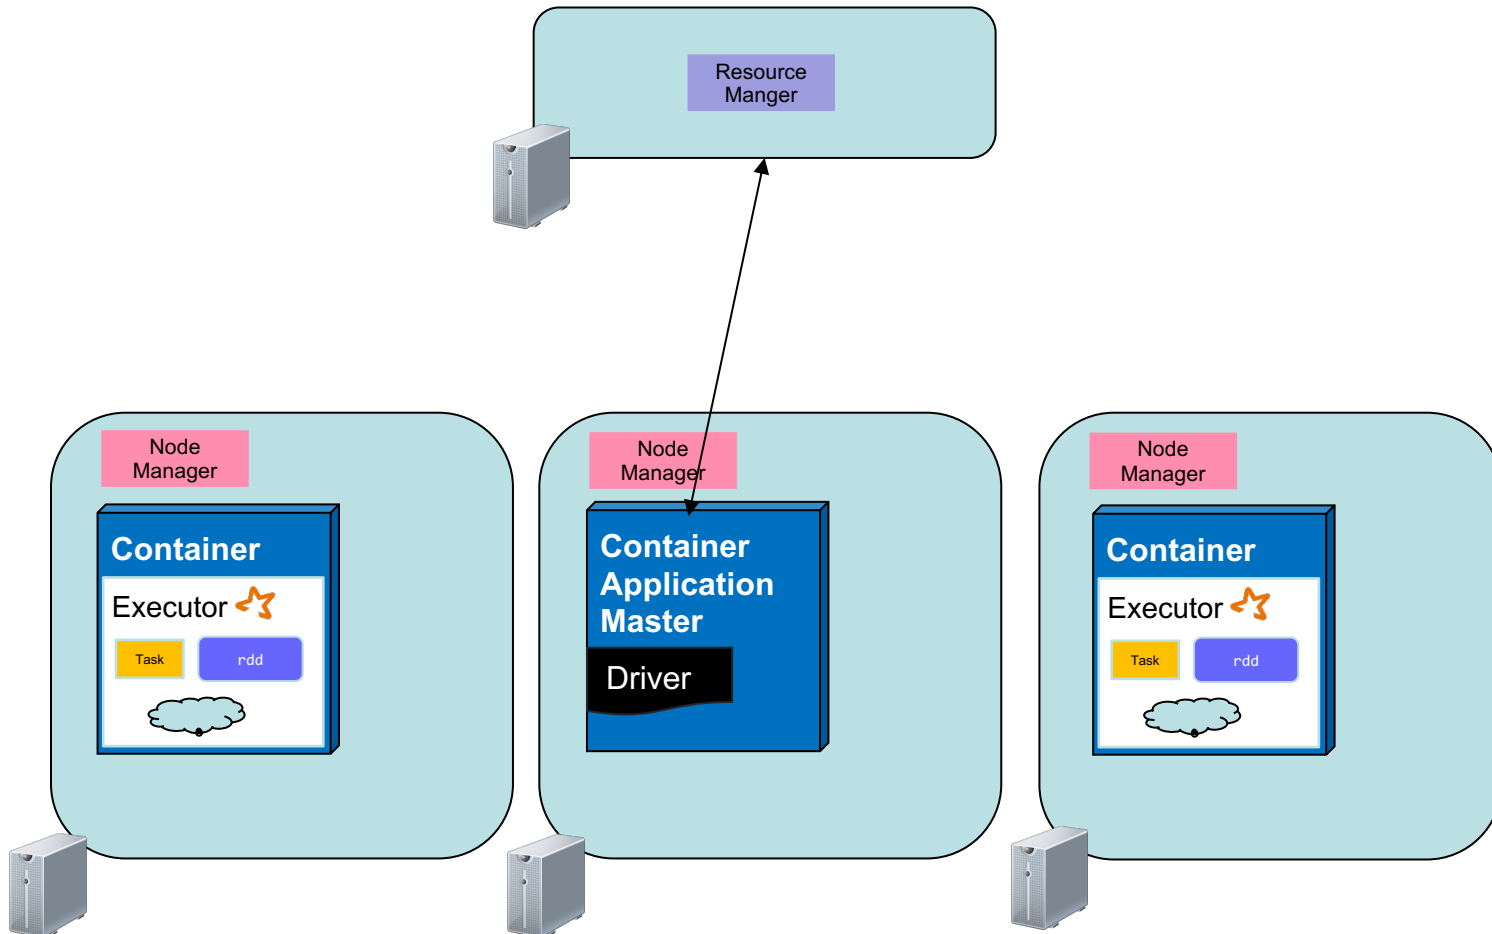
# Standalone “Cluster” Mode



`./bin/spark-submit --name "SecondApp" --master spark://host:port myApp.jar`

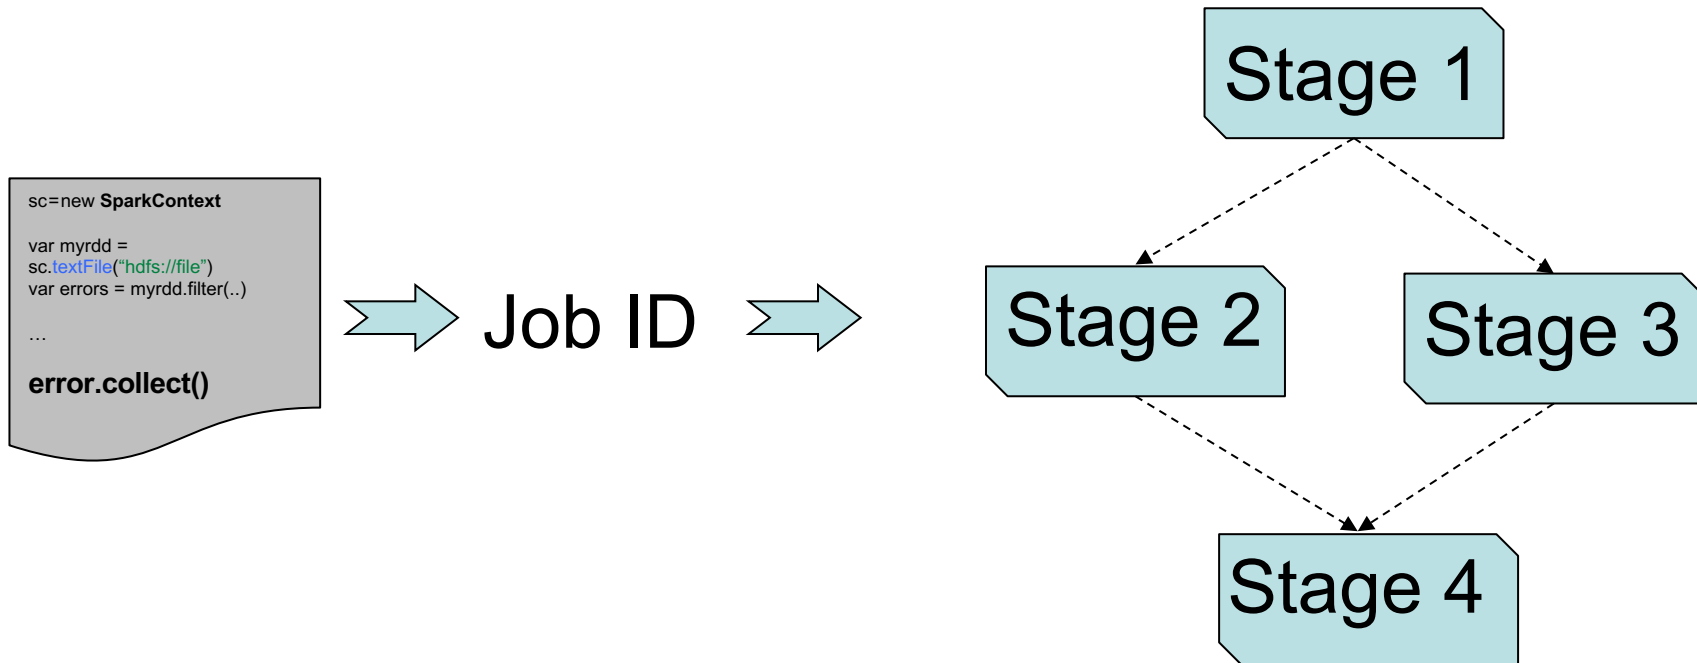


# YARN Mode

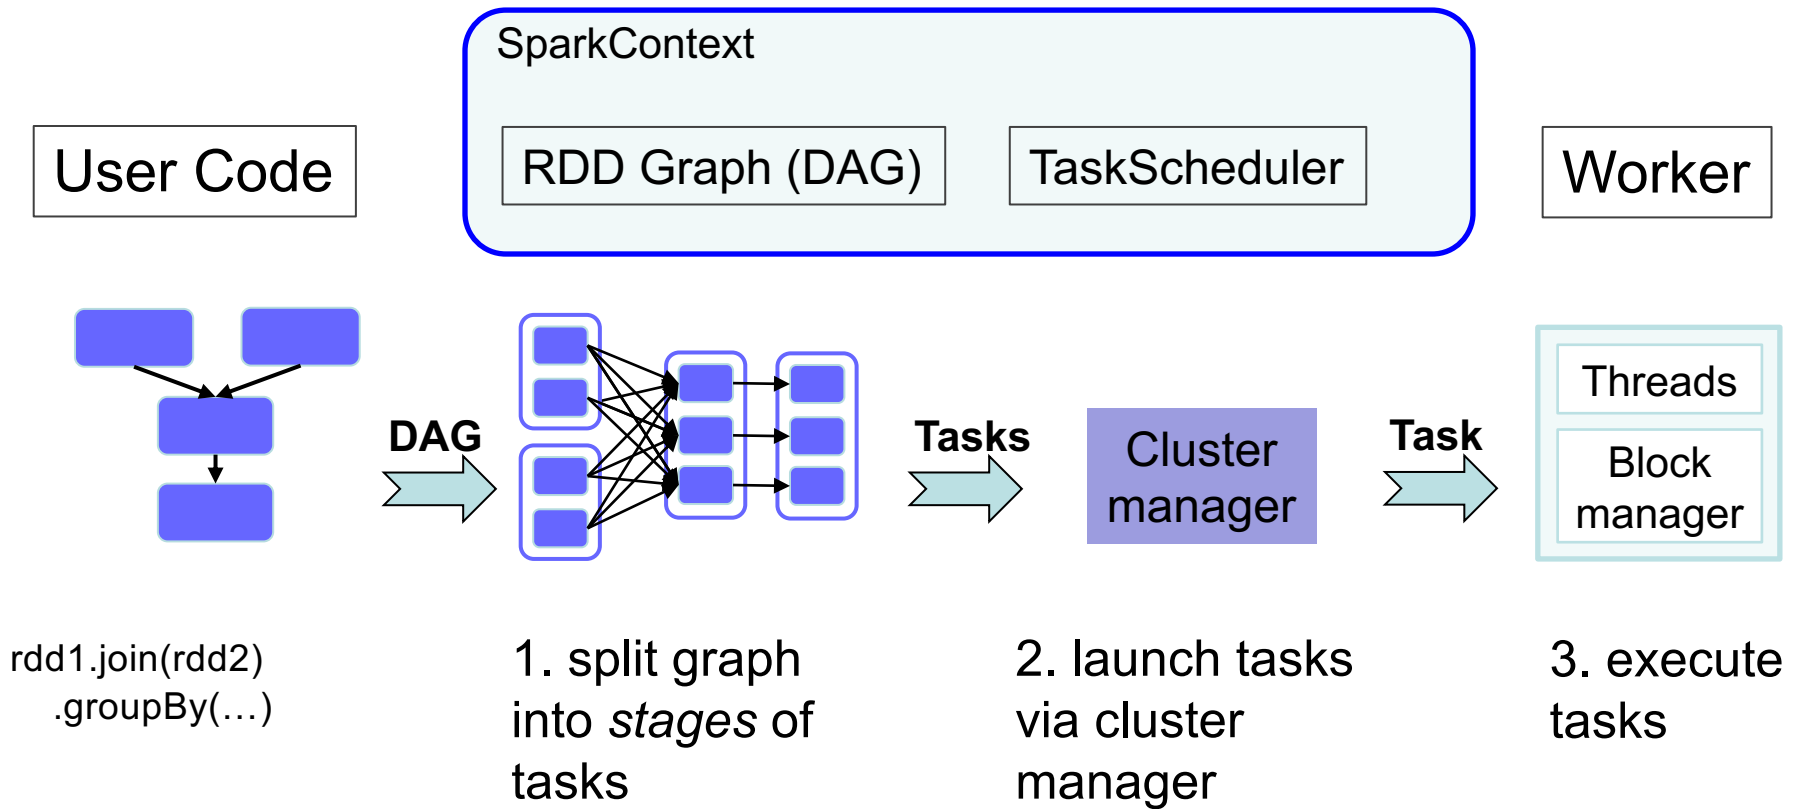


# Staged Execution

*Given a job, Spark generates the stages of execution*



# Scheduling Process



# Lineage

---

- One of the challenges in providing RDDs as an abstraction is choosing a representation for them that can track lineage across a wide range of transformations.
  - How to represent dependencies between RDDs?
- In practice, classify dependencies into two types
  - **narrow dependencies**, where each partition of the parent RDD is used by at most one partition of the child RDD
  - **wide dependencies**, where multiple child partitions may depend on it.

(Zaharia 2012)

# DAG Scheduling

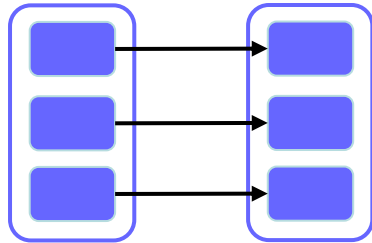
---

1. Spark creates an **operator graph** on the user code (RDD lineage)
2. When an **Action**, the operator graph is submitted to the DAG Scheduler.
3. The DAG Scheduler breaks the lineage into stages **based on the presence of wide dependencies**.
  - Each stage consists of a set of tasks that can be executed together on the same set of input data. 💡 Spark optimizes the execution plan to minimize the number of shuffle operations required.
4. The stages are then passed on to the **Task Scheduler**, which launches tasks through the cluster manager.
5. The workers execute the tasks on the worker node. Spark coordinates the execution of tasks across the executors to ensure fault-tolerance and efficient resource utilization

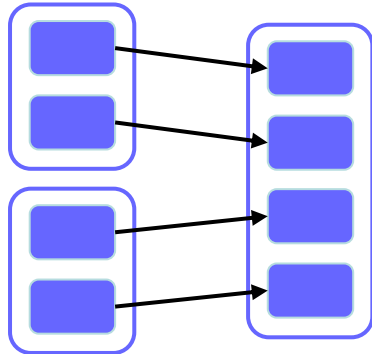
# Dependency Types

---

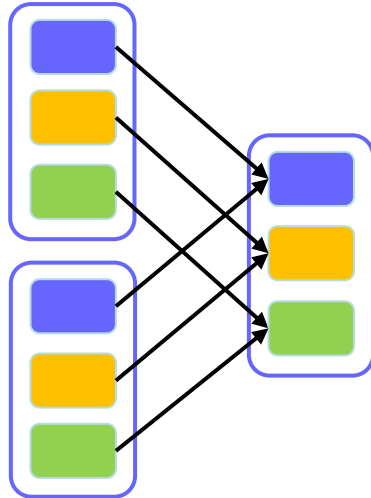
“Narrow” dependencies:



map, filter

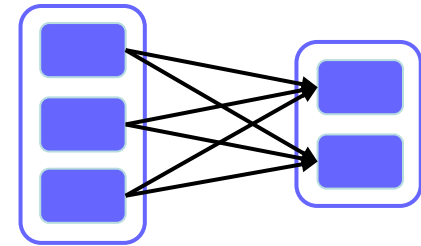


union

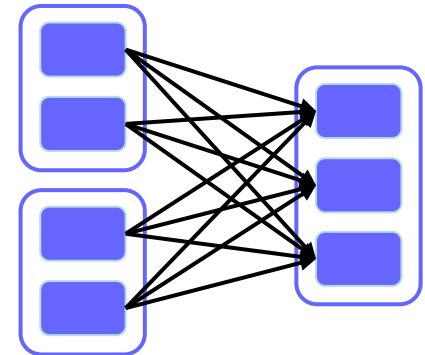


join with inputs  
co-partitioned

“Wide” (shuffle) dependencies :



groupByKey

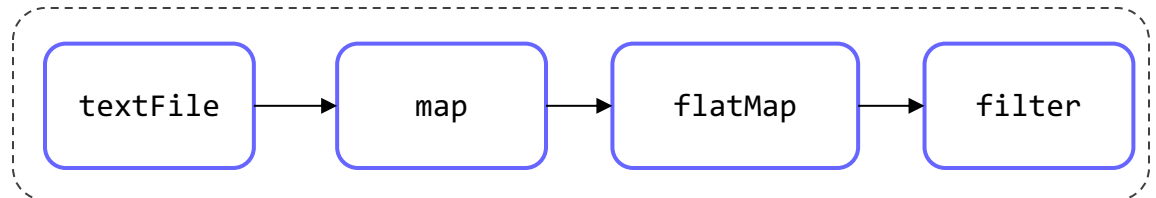


join with inputs not  
co-partitioned

# How Many Stages?

---

```
var a = sc.textFile("someFile.txt")  
  .map(mapFunc)  
  .flatMap(flatMapFunc)  
  .filter(filterFunc)  
  .count()
```



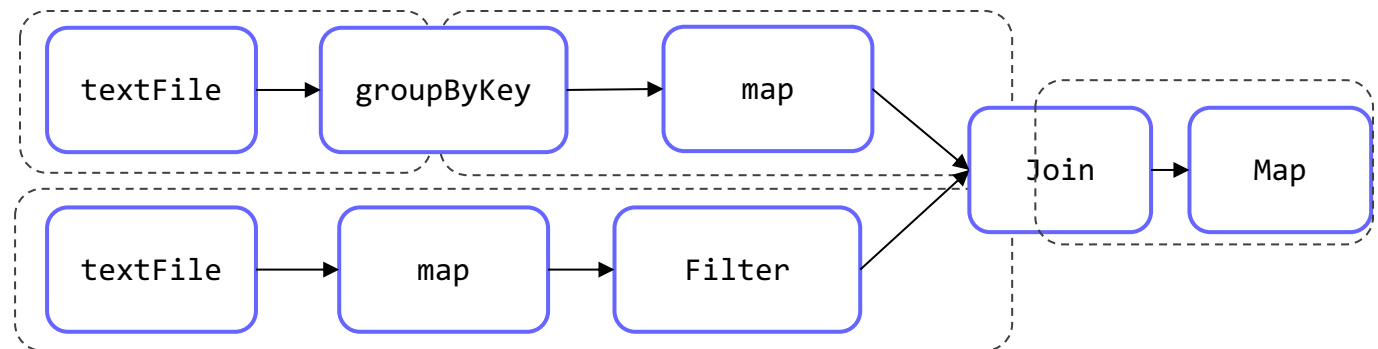
# How Many Stages?

---

```
var s = sc.textFile("sales")
```

```
var l = sc.textFile("locations")  
.groupByKey()  
.map()
```

```
s.map()  
.filter()  
.join(l)  
.map()  
.collect()
```





# Summary

---

- Spark is used for efficient distributed data processing:
  - In-memory
  - Lazy execution
  - MapReduce principles
  - Graph of executions
- RDD
  - A distributed data structure in Spark.
  - It defines how a collection (e.g., a list) is distributed on a cluster of machines.
- Getting Started with RDDs:
  - [Programming Guide](#)